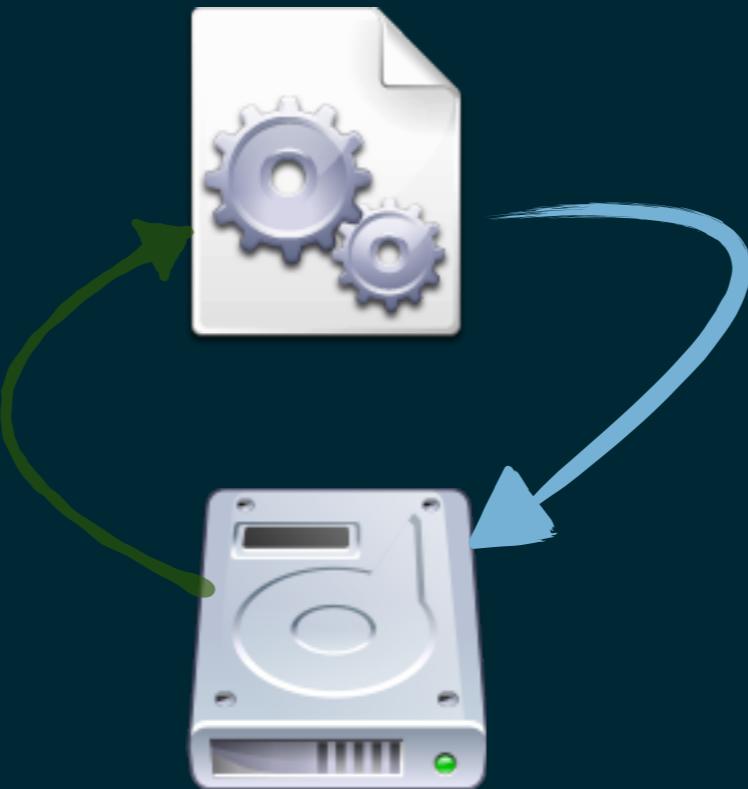




Wenn's mal wieder länger dauert
auf der Suche nach dem Bottleneck

Applikation



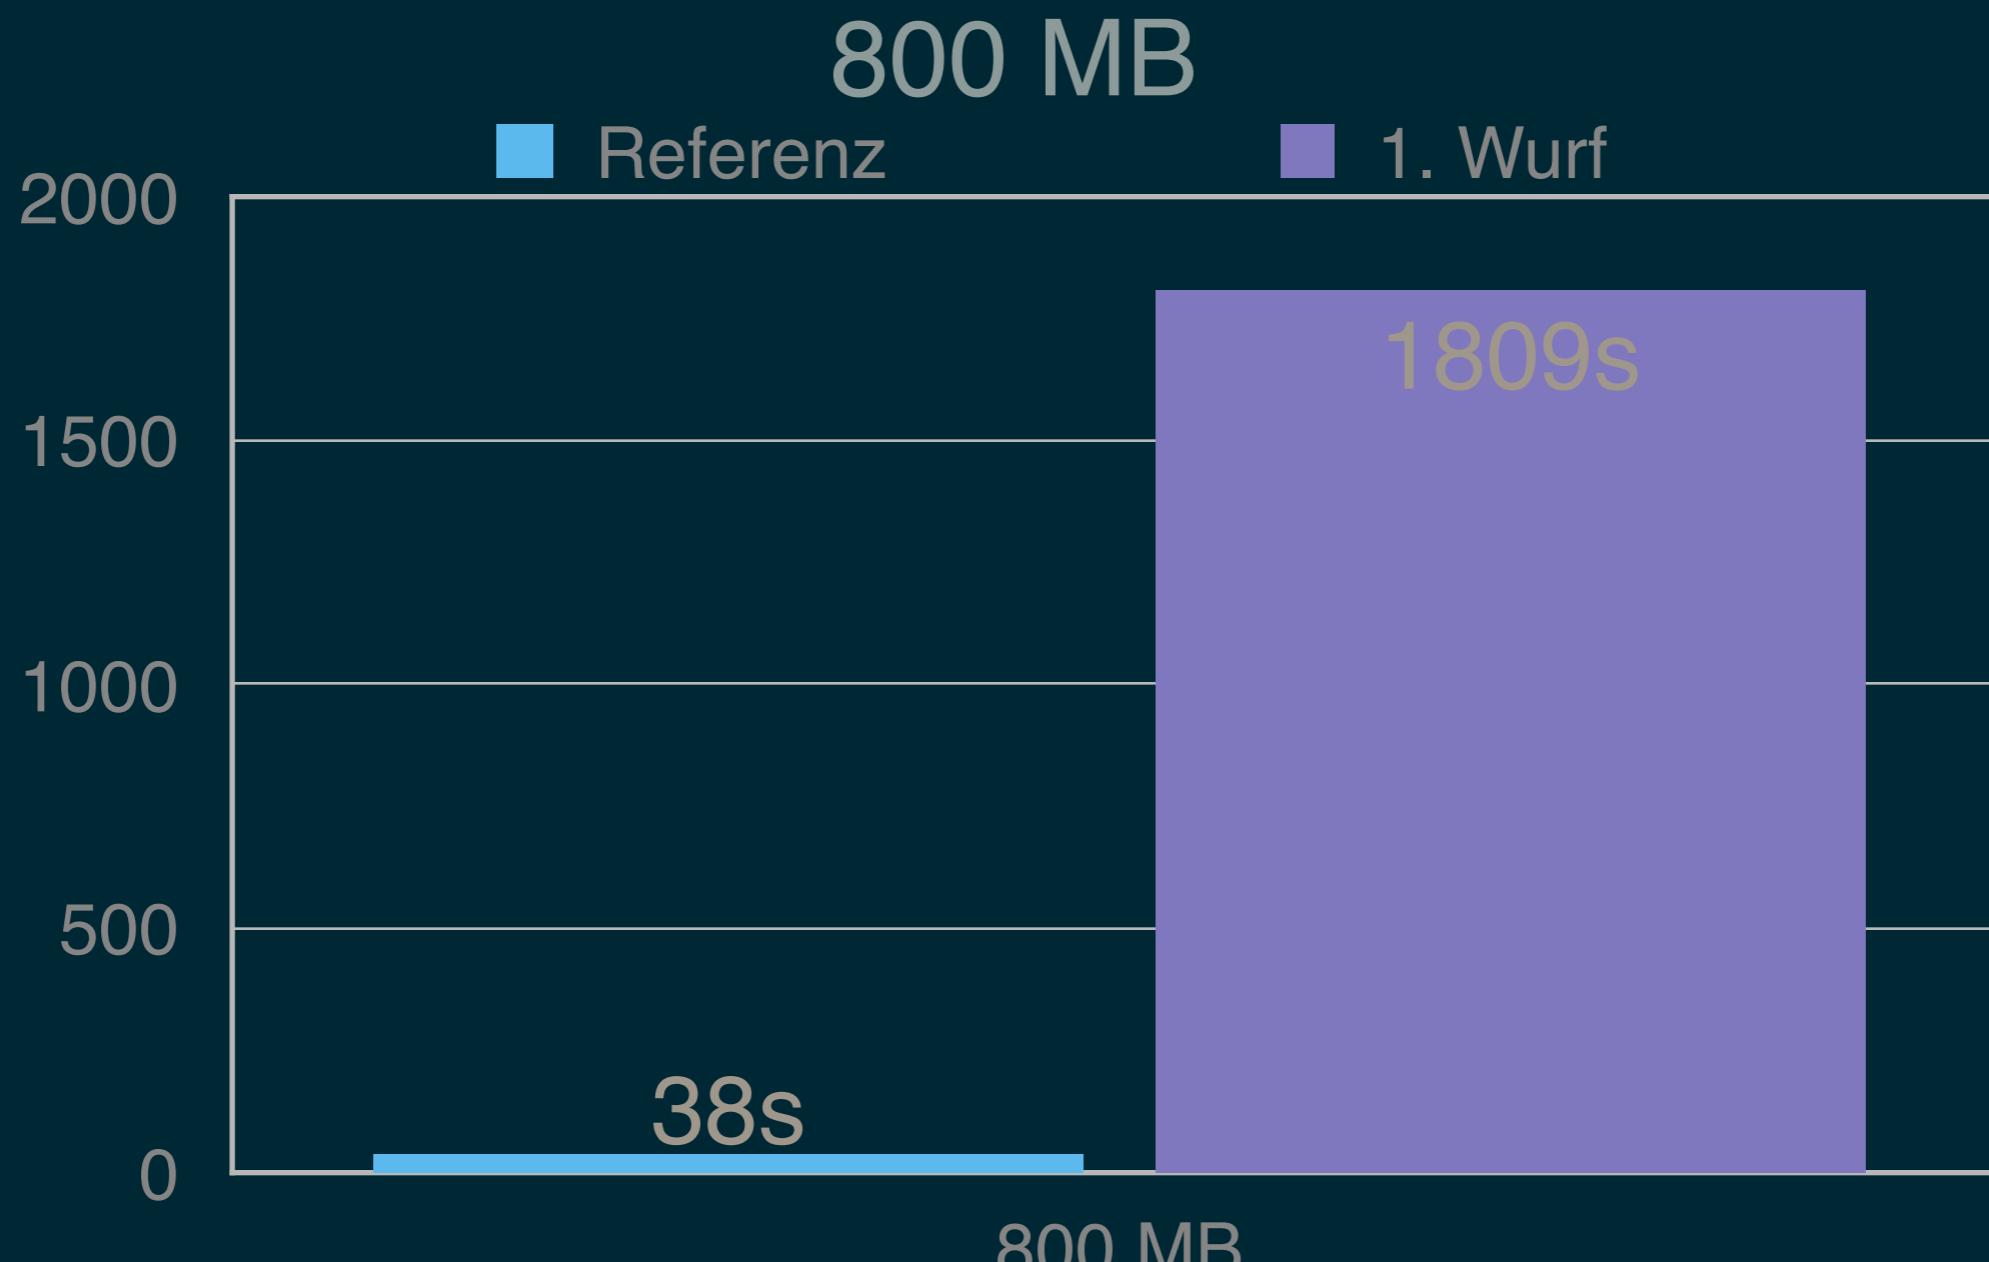
Datei im Streaming Verfahren verschlüsseln

Erwartungshaltung

Wenn ein optimiertes C
Programm 40 Sekunden braucht,
wie lange braucht dann ein
optimiertes Java Programm?



Laufzeit/Datenmenge



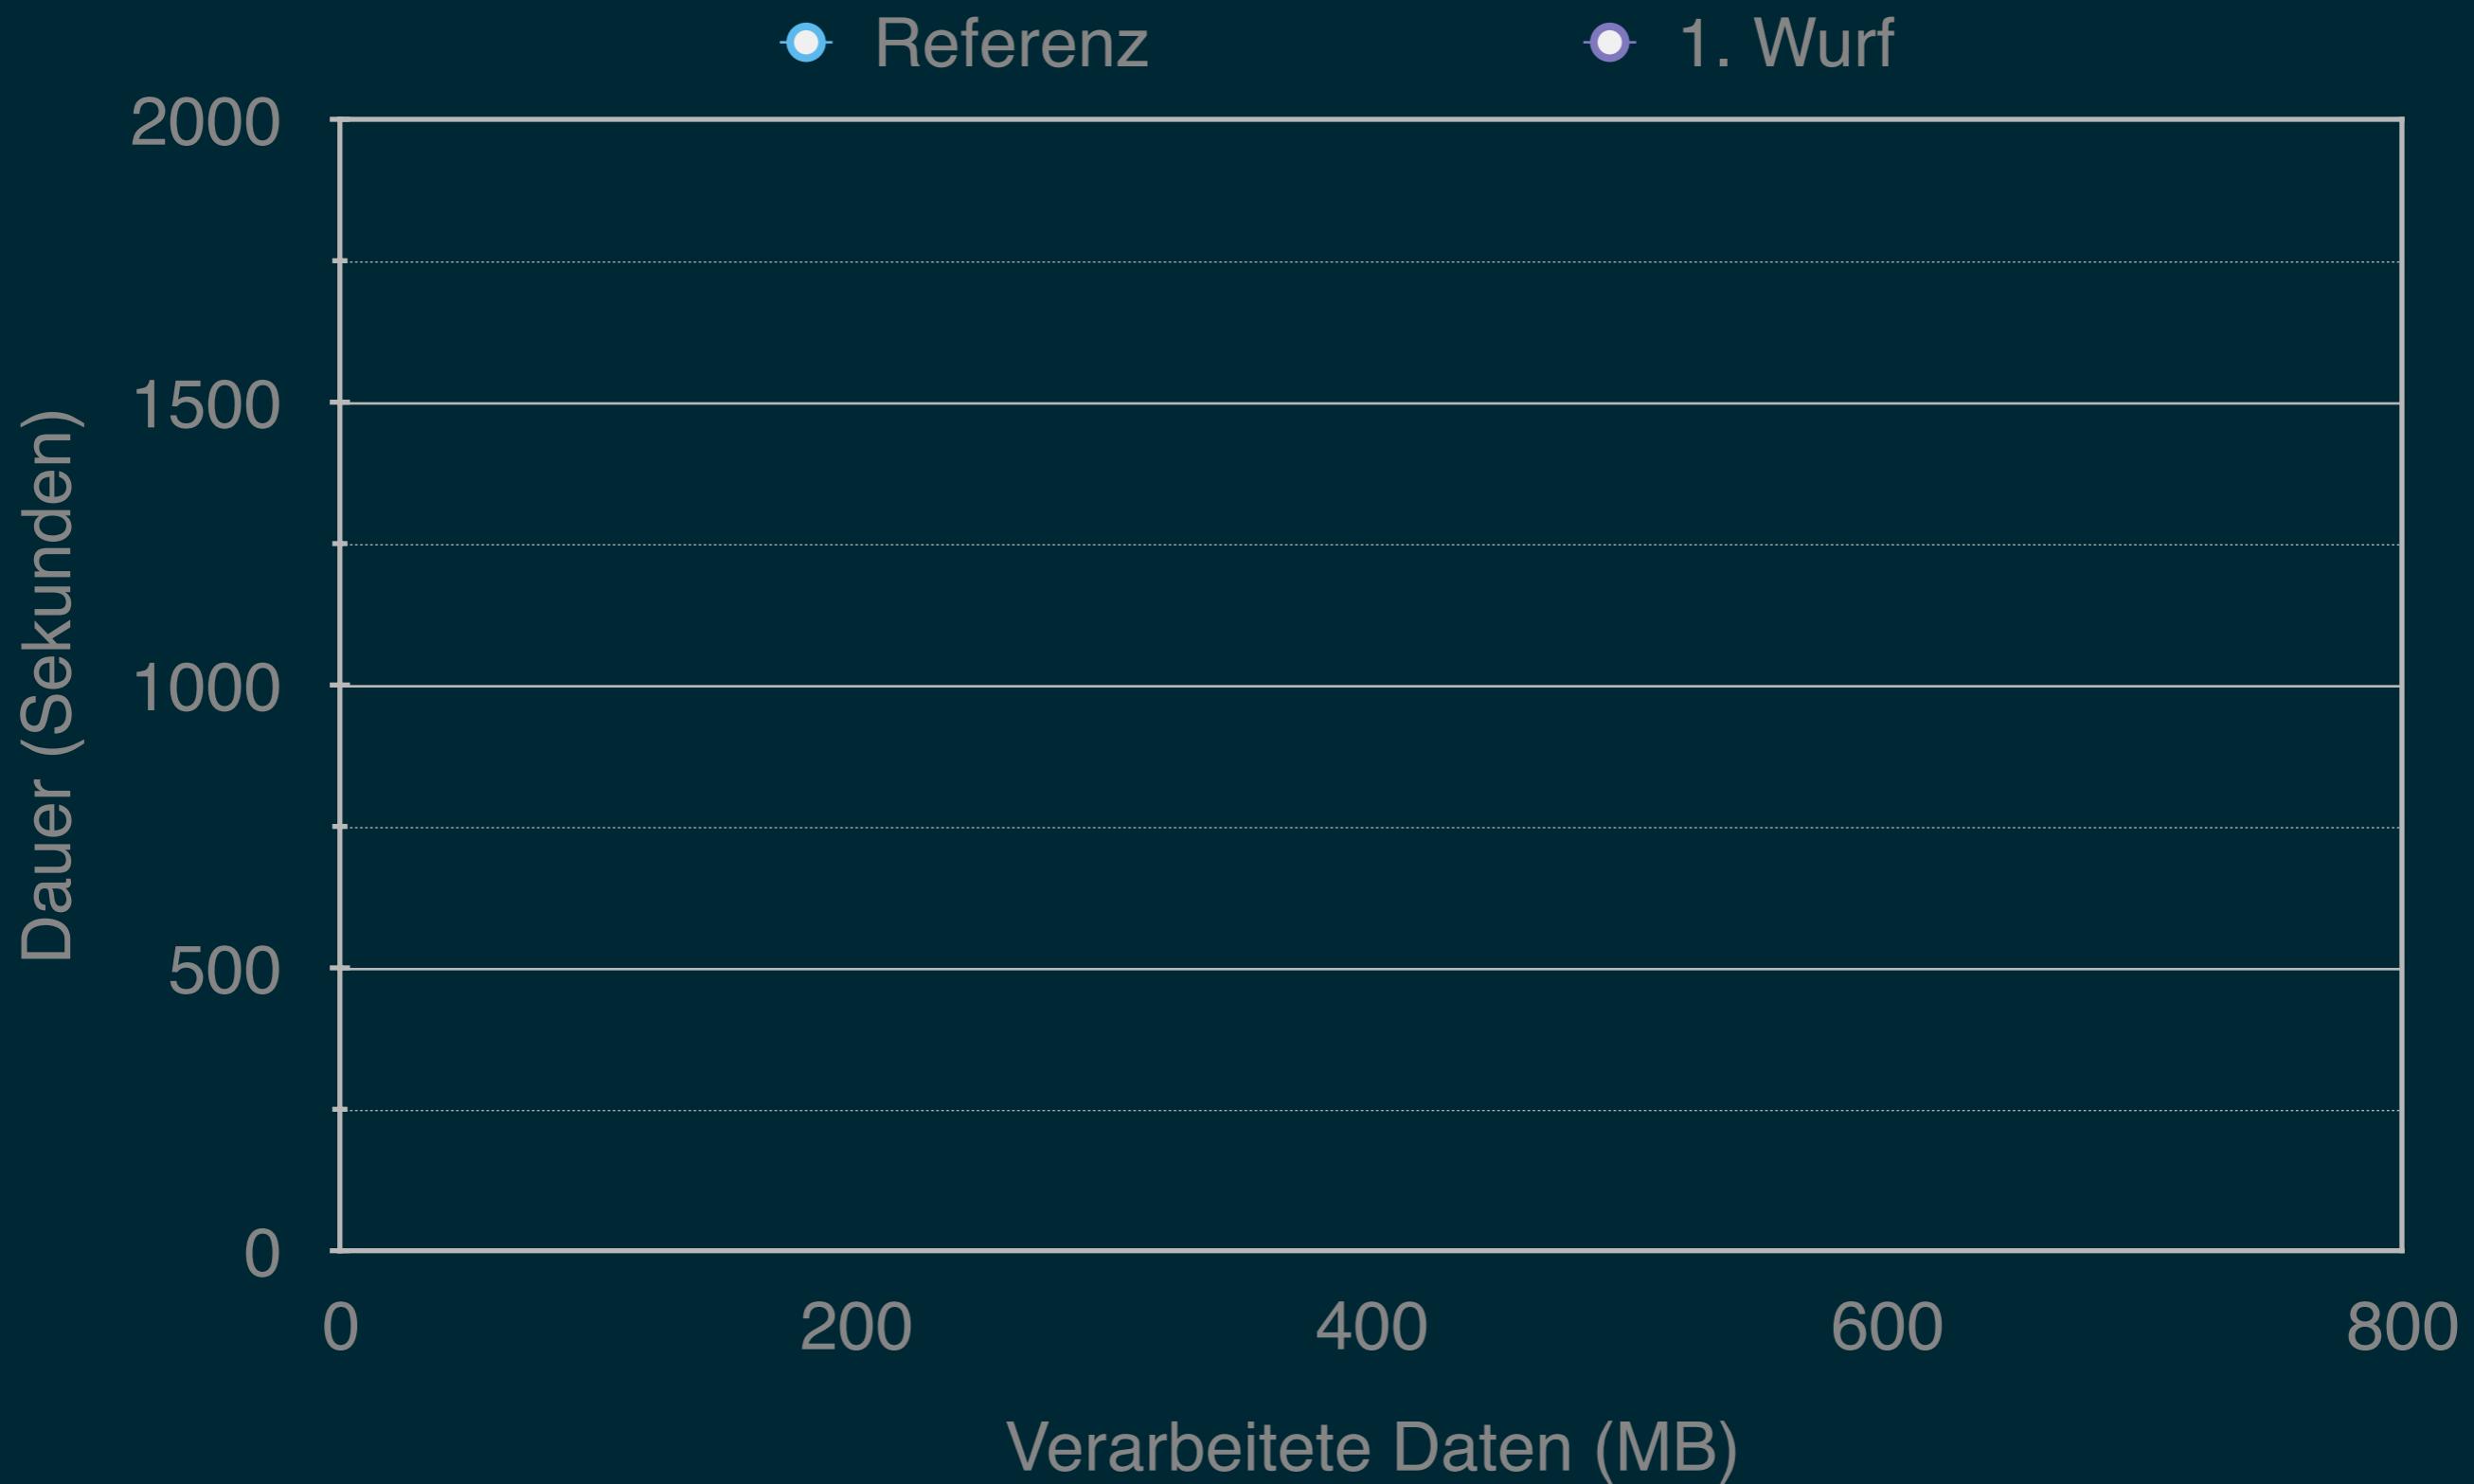
~ Faktor 47 zwischen Referenz und 1. Wurf

Happy?

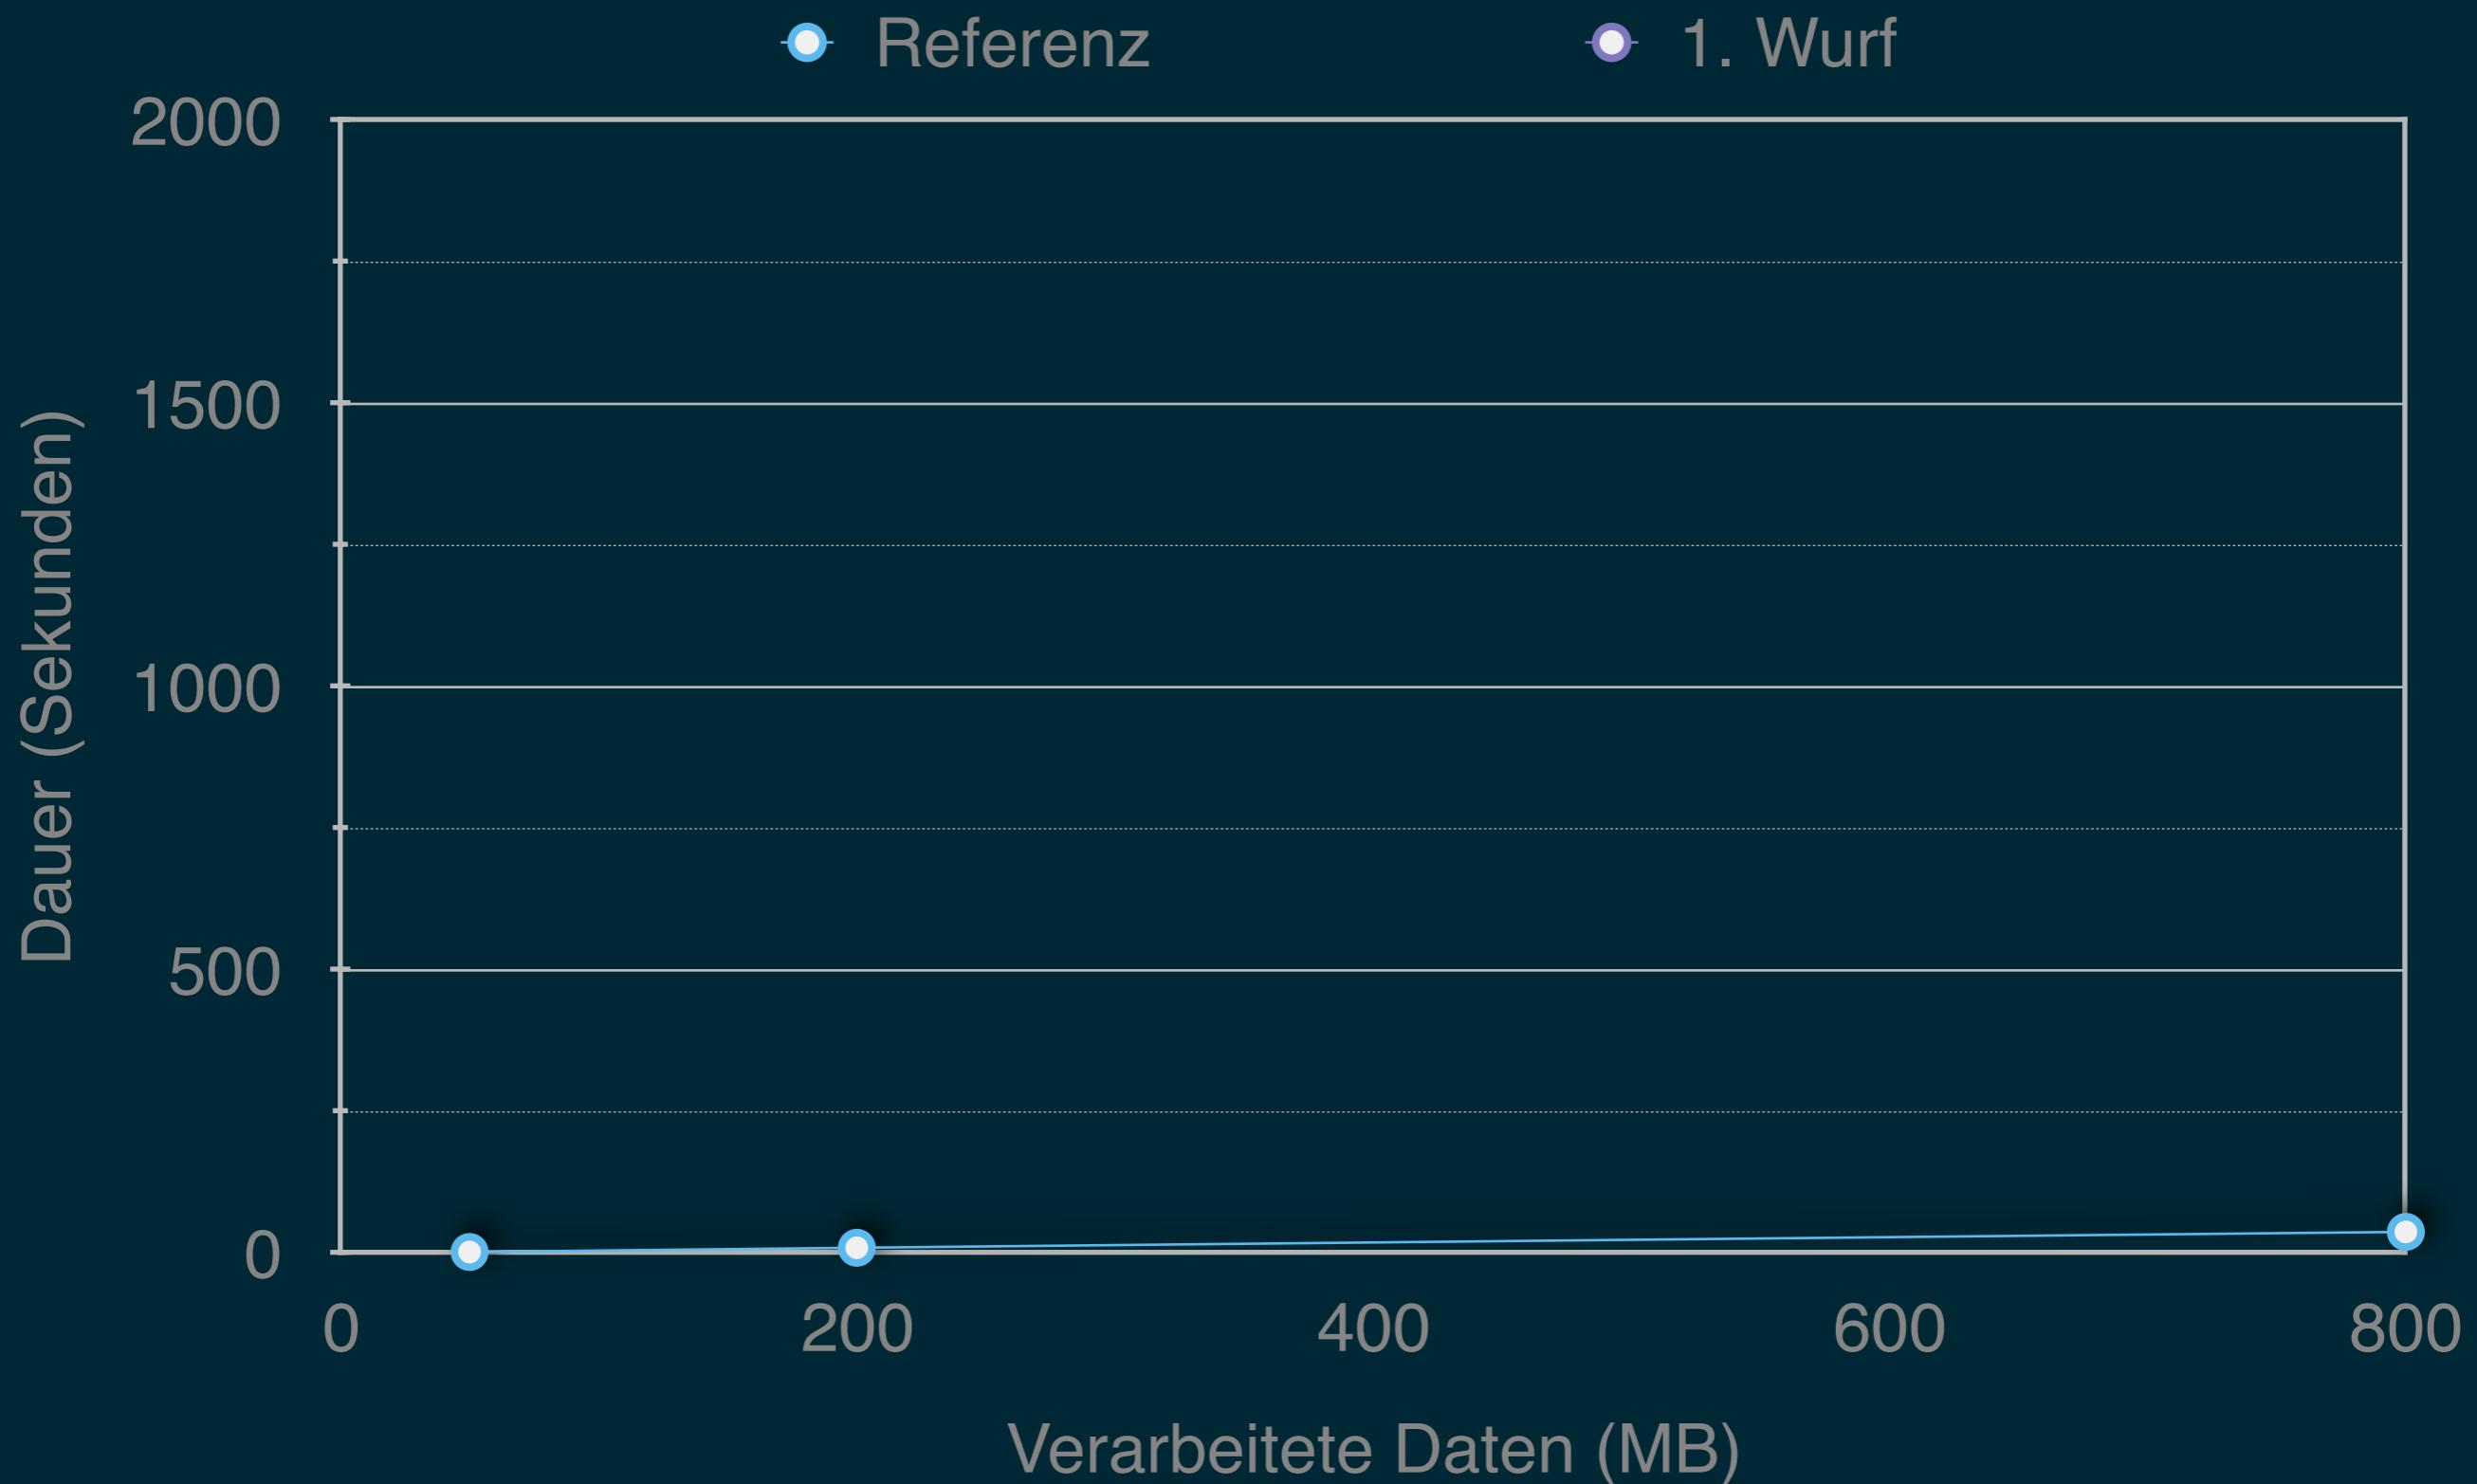


Laufzeit/Datenmenge

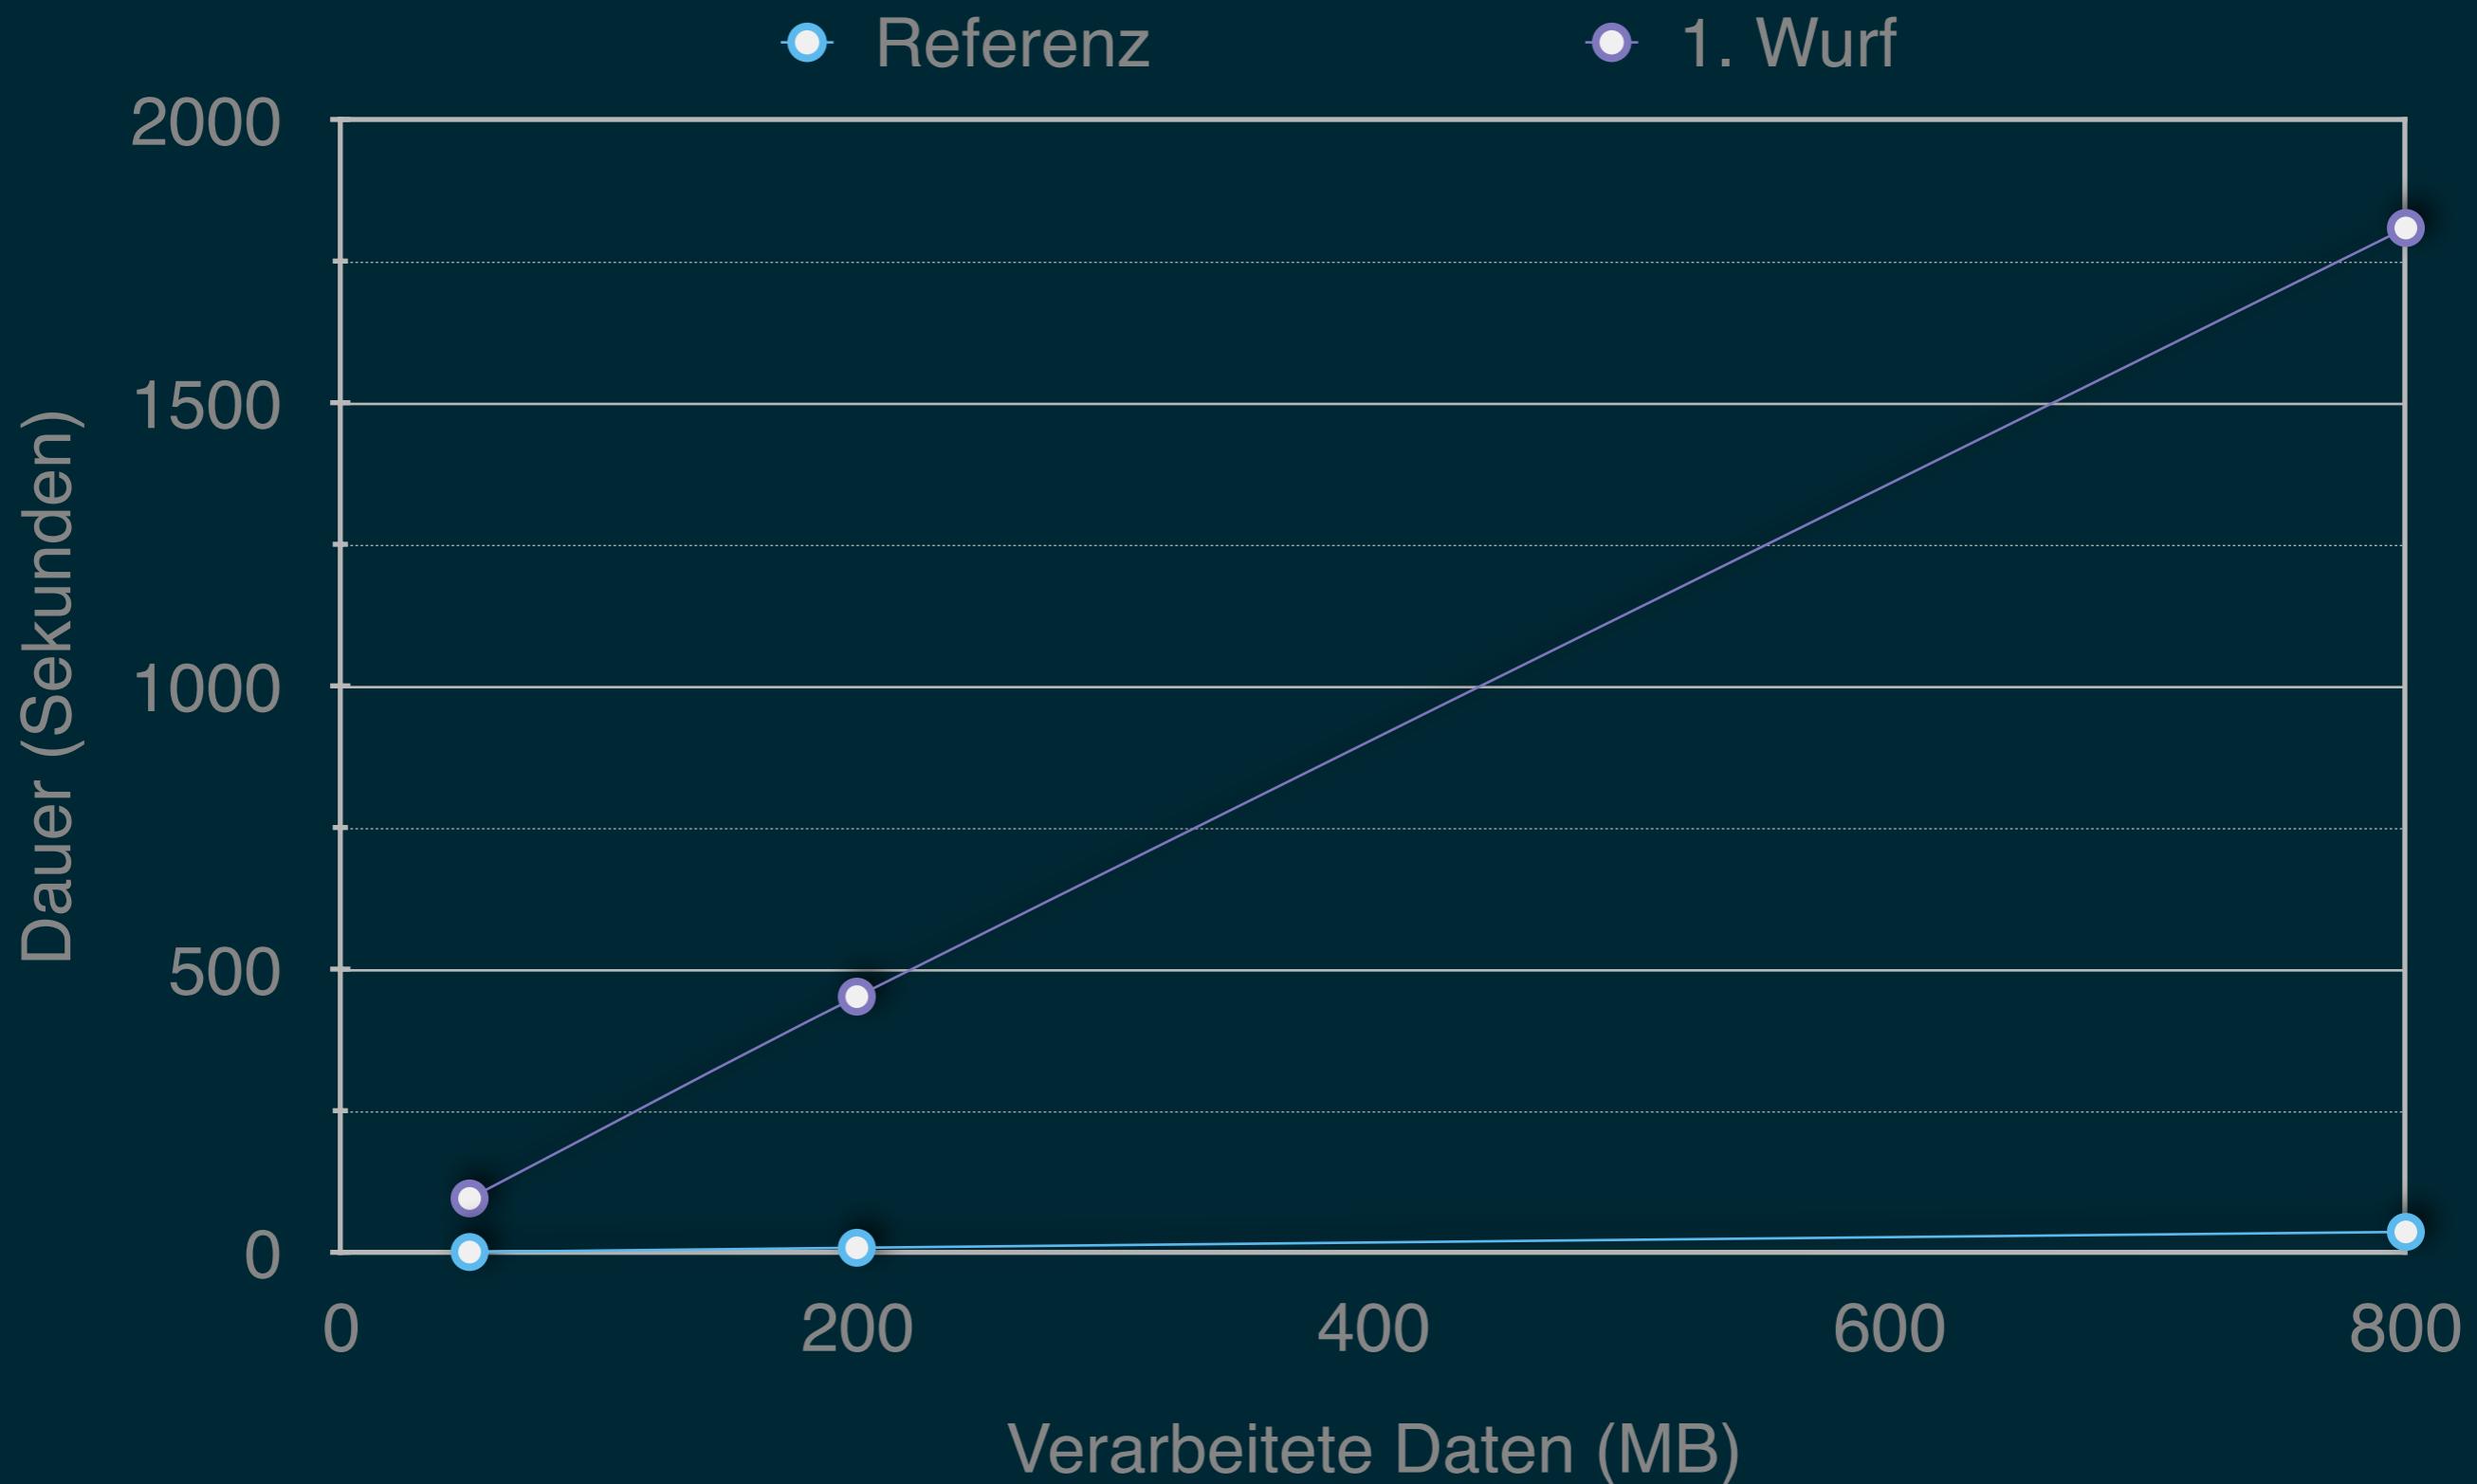
Laufzeit/Datenmenge



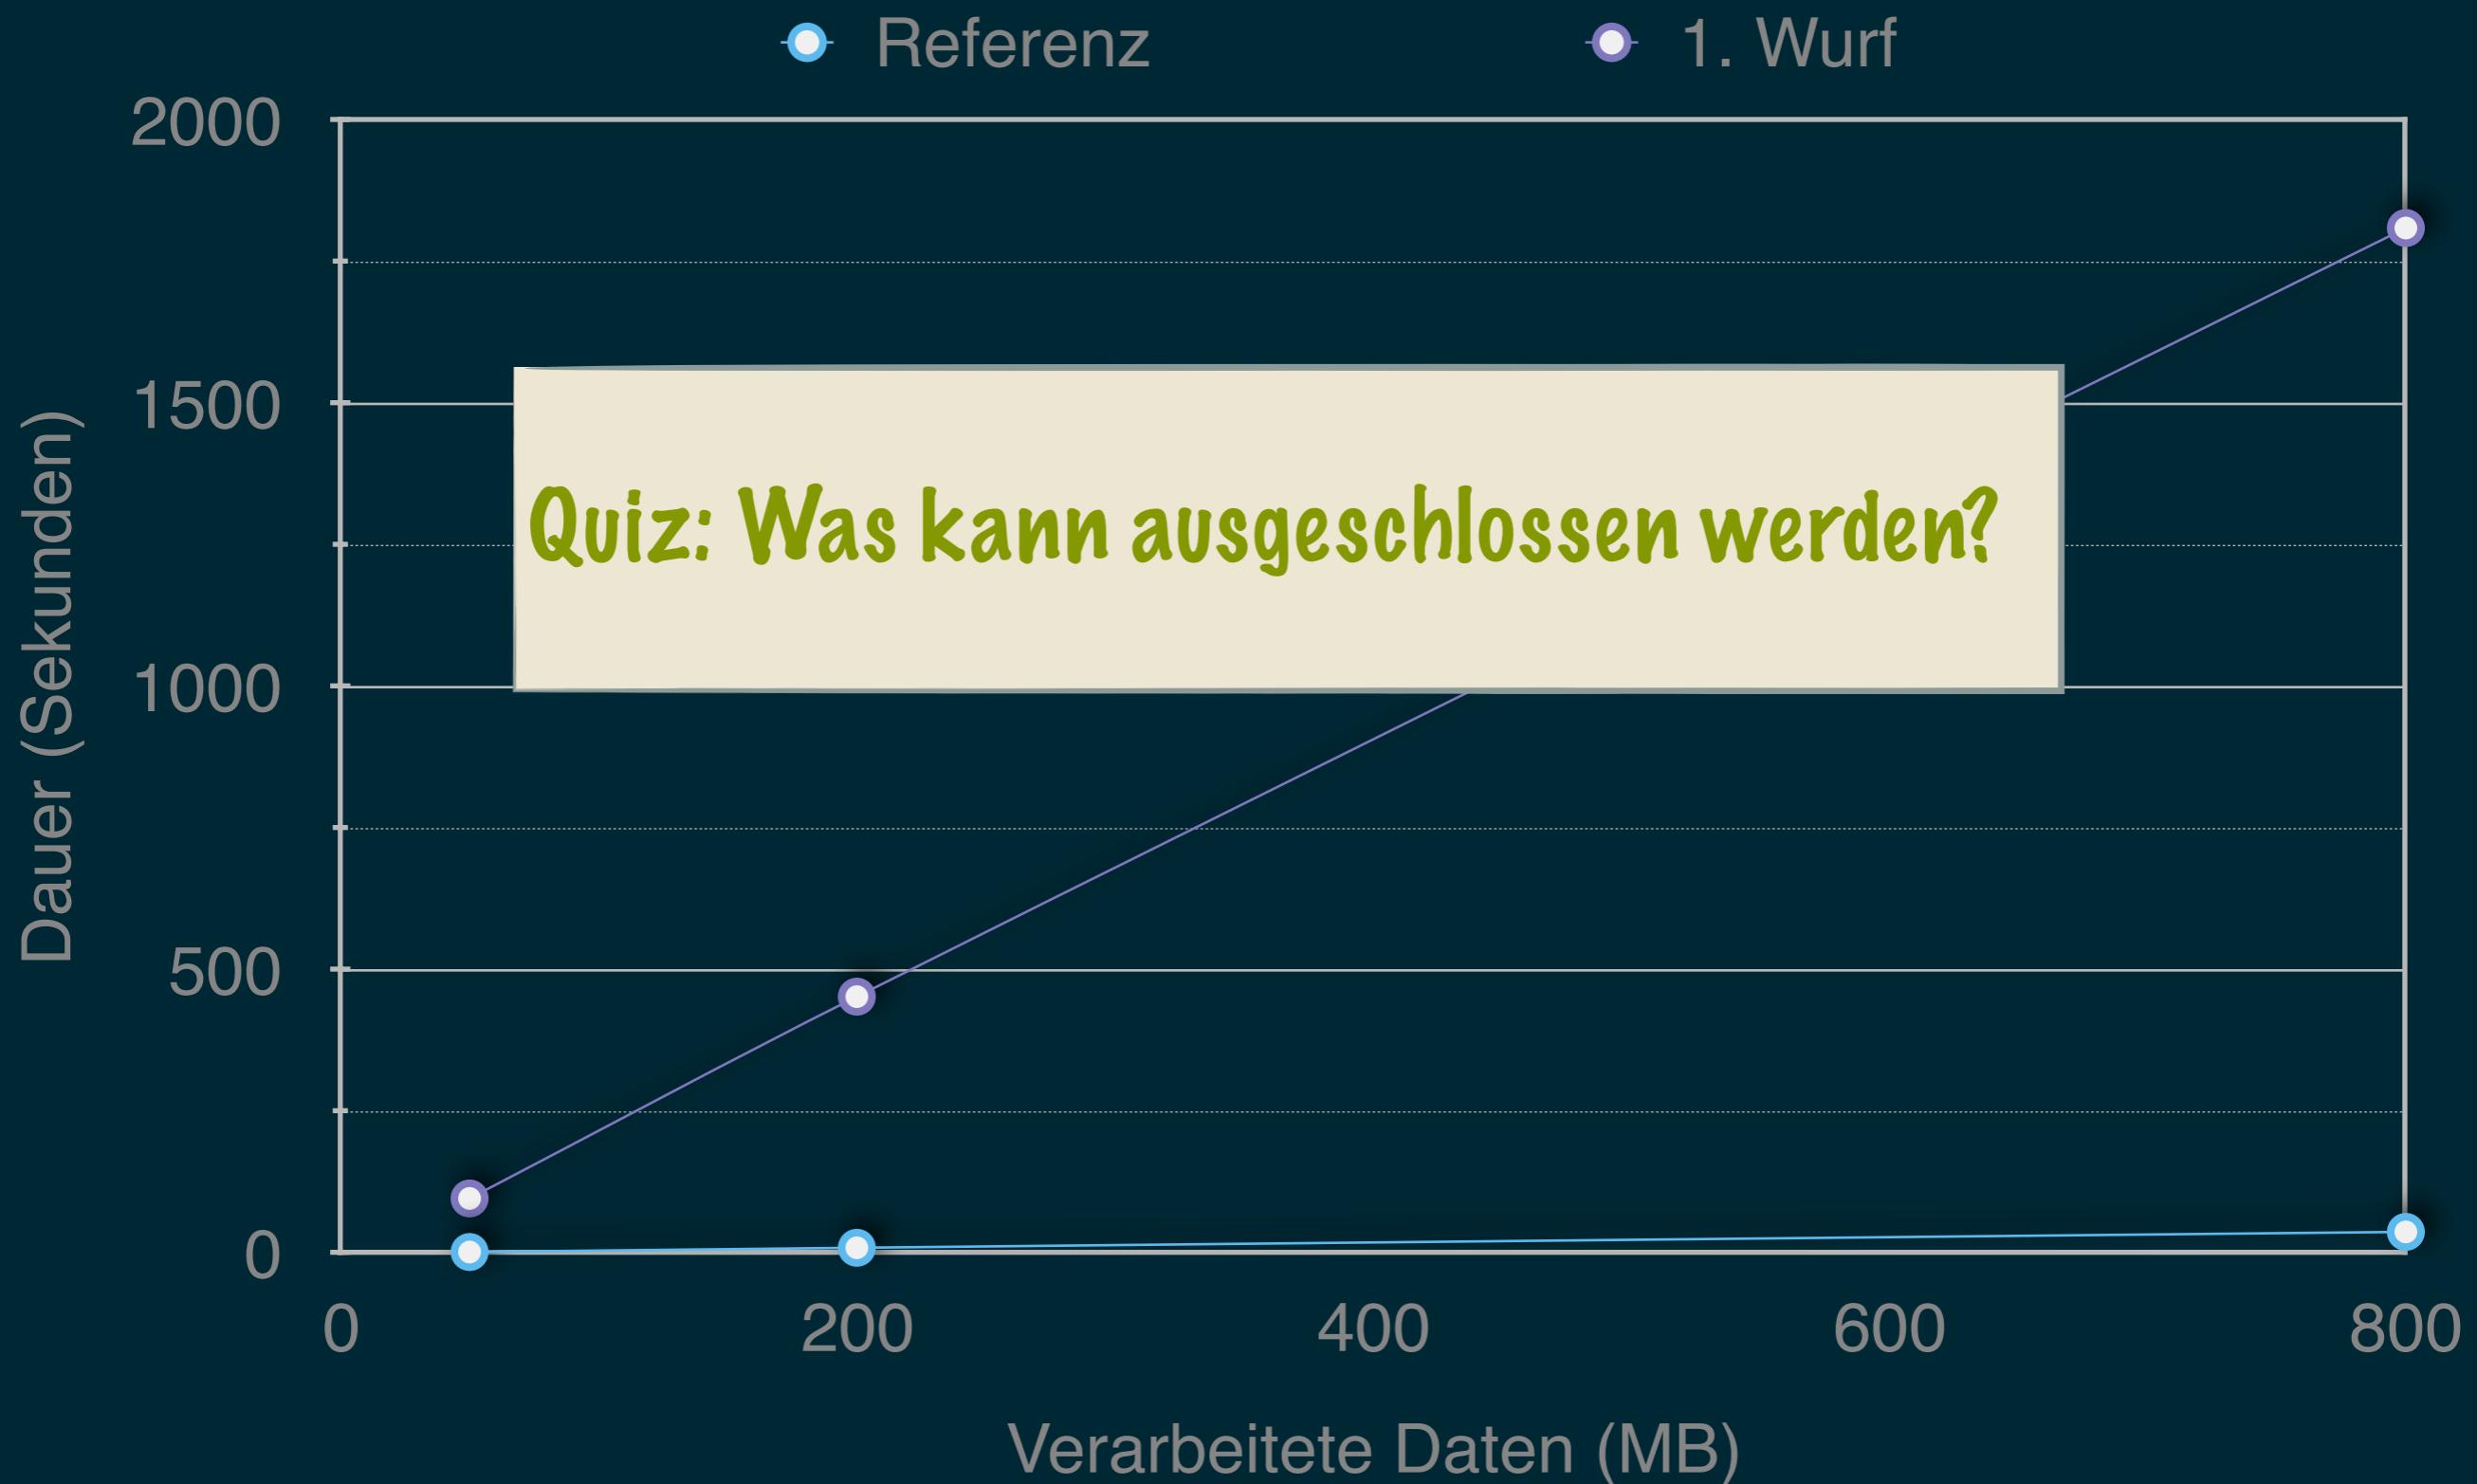
Laufzeit/Datenmenge



Laufzeit/Datenmenge



Laufzeit/Datenmenge



Und jetzt?



Und jetzt?

Das Problem einkreisen!



Was tut der Rechner?

- * Rechnen (CPU)*
- * I/O (Platte, Netz)
- * Blockieren (z.B. Mutex)



Bottleneck?

- * Nicht Thema:
Multi-Core, CPU Caches, TLB thrashing, ...
Effiziente Algorithmen

Was tut der Rechner?

Begriffe

- * *Ressource*: CPU, Disk, Network, ...
- * *Utilisation*: Wie viel Zeit (in %) ist eine Ressource beschäftigt?
- * *Saturation*: Wie viel Arbeit “bleibt liegen”?
- * *Errors*: Wie viele Fehler treten auf?

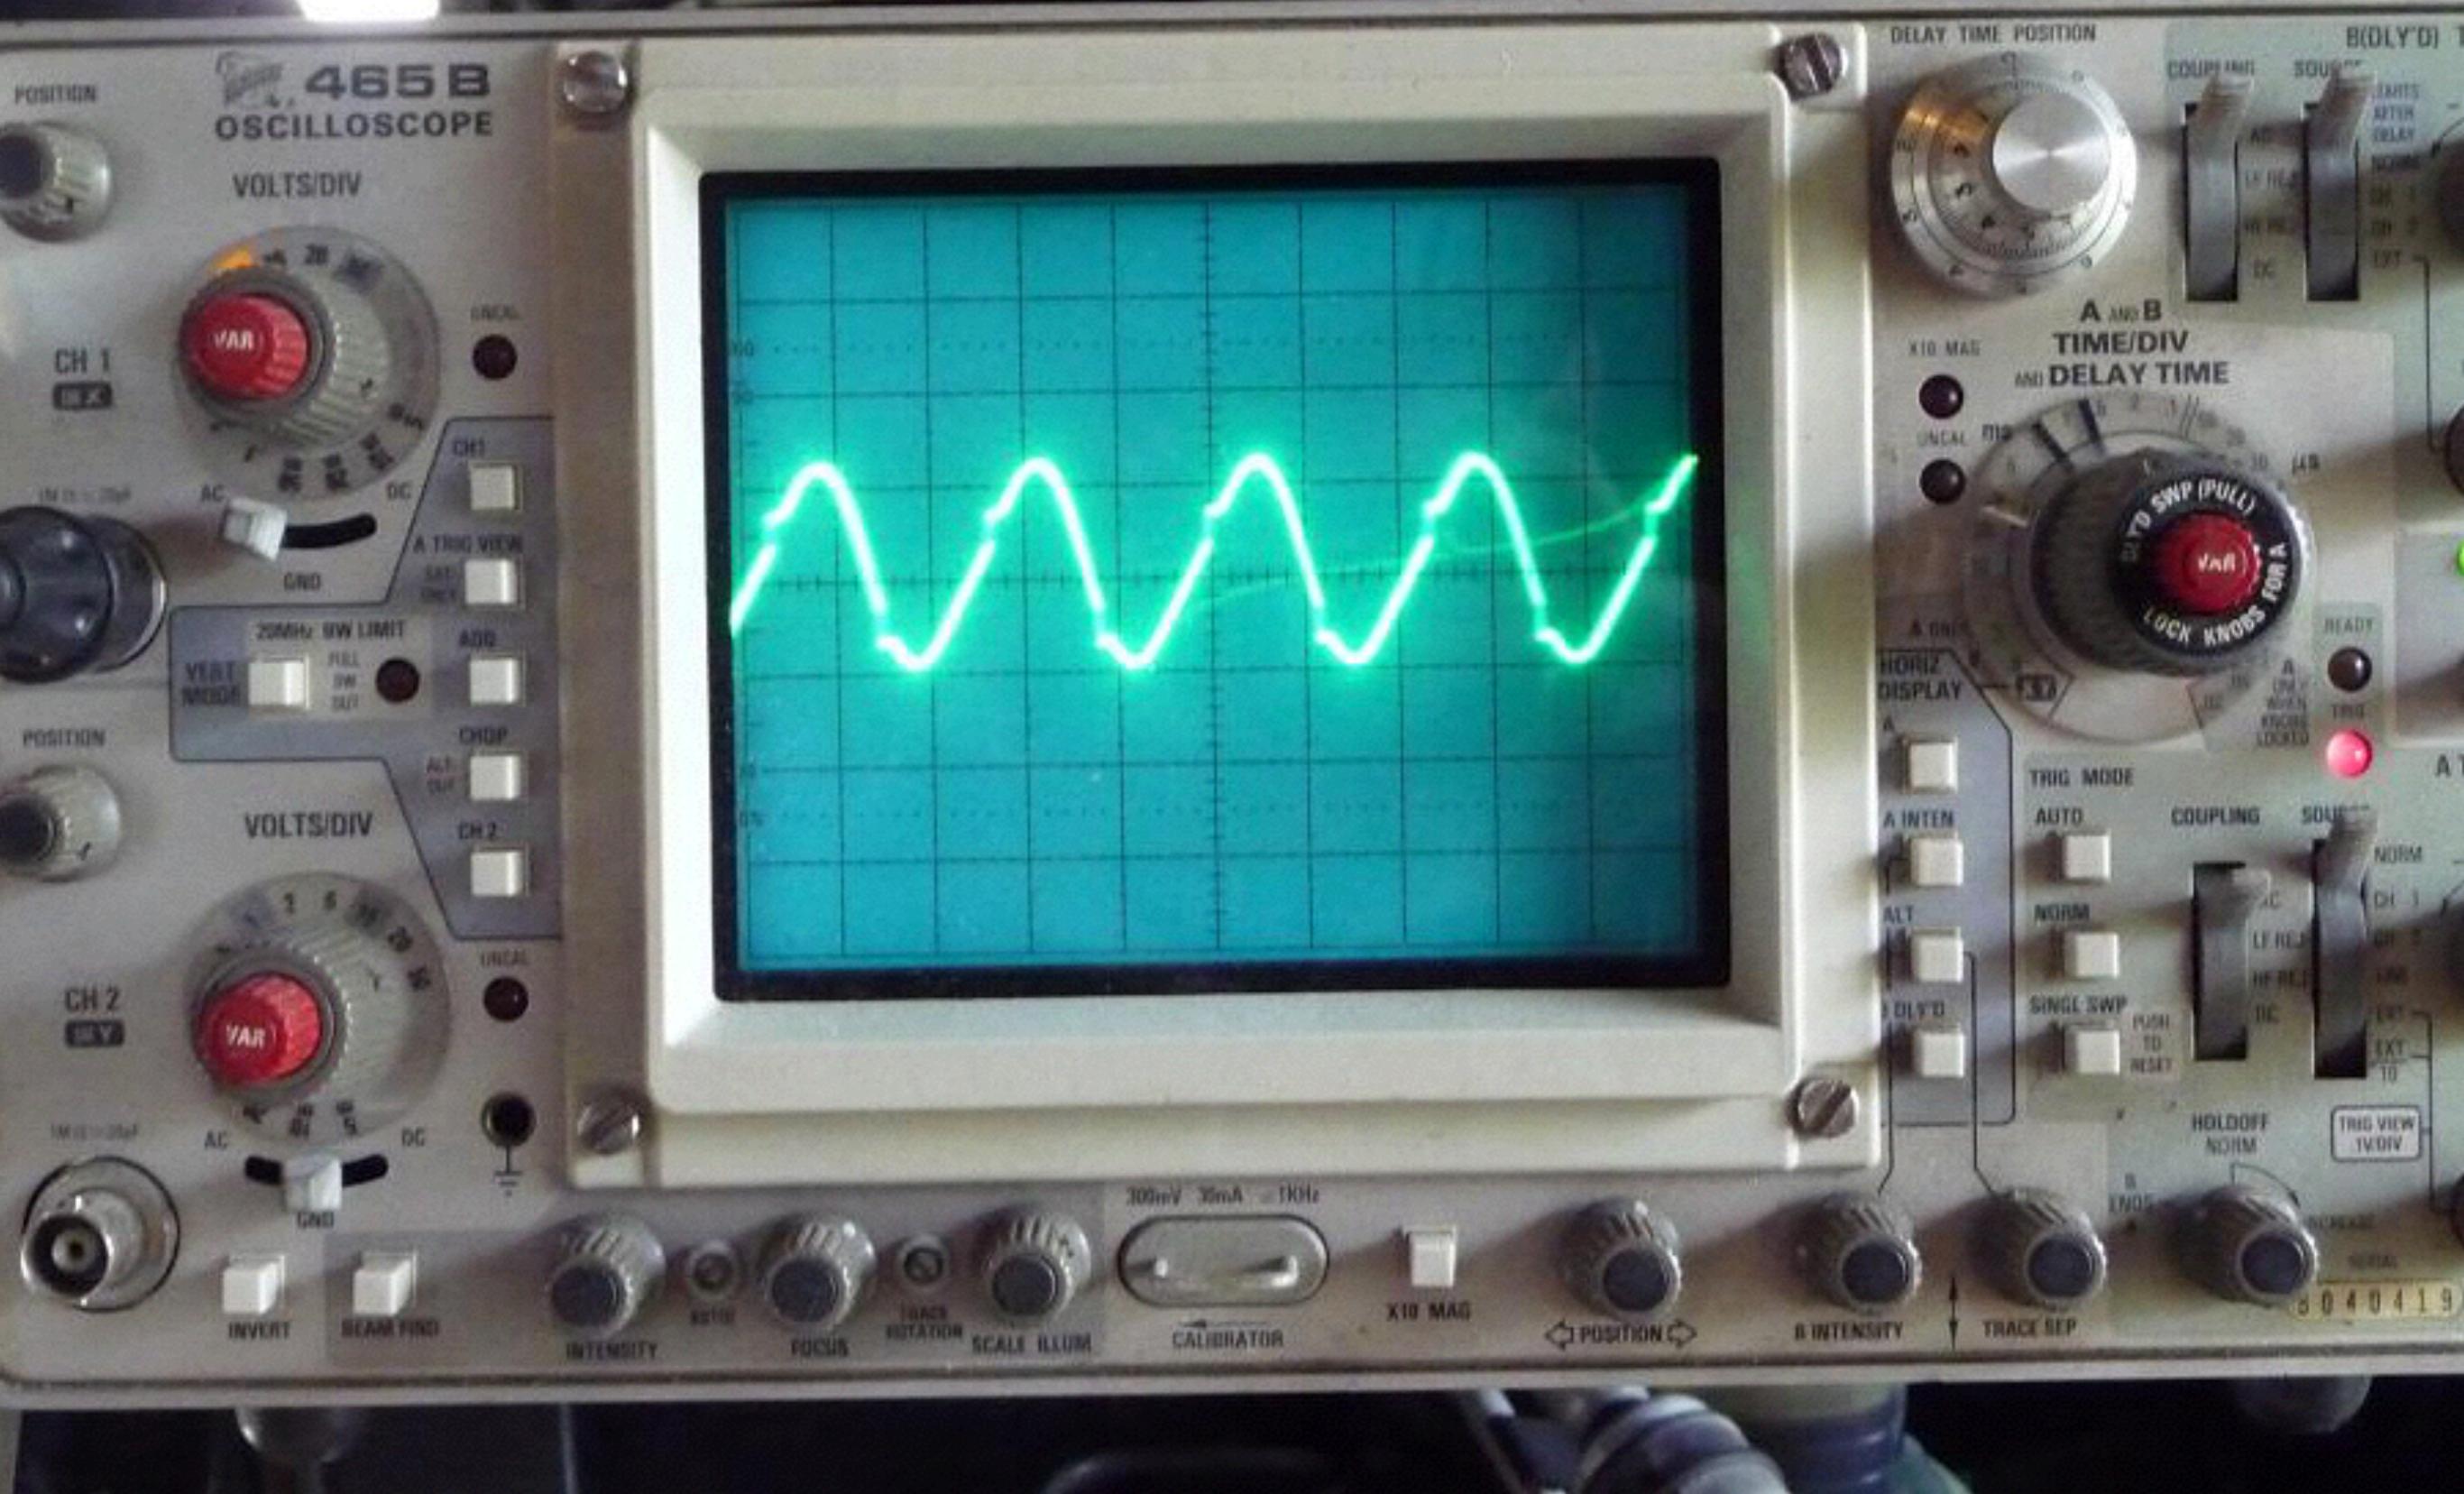


Was soll er tun?

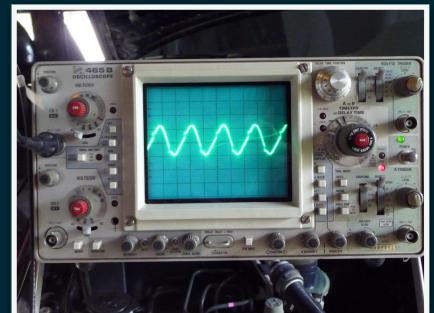
oder: Wo ist der Bottleneck?

- * *Für jede Ressource:*
 - * *Utilisation:* beschäftigt! *
 - * *Saturation:* Es bleibt keine Arbeit liegen!
 - * *Errors:* Es treten keine Fehler auf!
- * Je nach System/Messung: Utilisation sollte ~80% sein
Durchsatz vs. Latenz: beachten!

Und jetzt? Messen!



Disk messen



- * *Utilisation:* iostat (busy time, IOPS)
- * *Saturation:* iostat (wait time, IOPS), vmstat (b)
- * *Errors:* iostat

Disk



```
$ iostat -x -p dm-3 60
```

rrqm/s	wrqm/s	r/s	w/s	rkB/s	wkB/s	avgrq-sz	avgqu-sz	await	r_await	w_await	svctm	%util
0.00	0.00	0.17	6.50	0.67	141.07	42.52	0.05	6.91	4.00	6.98	0.70	0.47

```
$ iostat -x -p dm-3 60
```

[...]	r/s	w/s	rkB/s	wkB/s	[...]	svctm	%util
[...]	0.17	6.50	0.67	141.07	[...]	0.70	0.47

Disk



```
$ iostat -x -p dm-3 60
```

rrqm/s	wrqm/s	r/s	w/s	rkB/s	wkB/s	avgrq-sz	avgqu-sz	await	r_await	w_await	svctm	%util
0.00	0.00	0.17	6.50	0.67	141.07	42.52	0.05	6.91	4.00	6.98	0.70	0.47

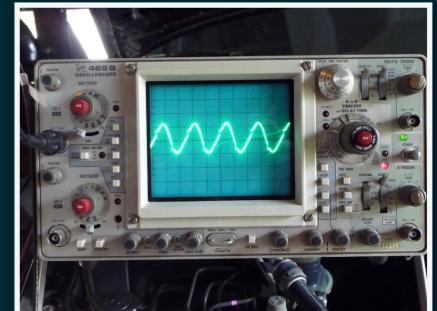
```
$ iostat -x -p dm-3 60
```

[...]	r/s	w/s	rkB/s	wkB/s	[...]	svctm	%util
[...]	0.17	6.50	0.67	141.07	[...]	0.70	0.47

* I/O: kein Problem

CPU

messen



- * *Utilisation:* time, top, vmstat, mpstat
- * *Saturation:* vmstat run queue > #cores
- * *Errors:* z.B. mcelog

CPU



```
$ time ./encrypt.sh [...]
```

```
[...]
```

```
Encryption took 452.47 s
```

```
real    7m32.730s  
user    0m16.597s  
sys     7m17.231s
```

```
$ mpstat -P ALL 60
```

CPU	%usr	%nice	%sys	%iowait	[...]	%idle
all	0.68	0.02	12.41	0.15	[...]	86.73
0	0.32	0.02	1.39	0.60	[...]	97.68
1	0.25	0.02	0.13	0.02	[...]	99.58
2	0.82	0.03	15.59	0.03	[...]	83.53
3	0.50	0.00	6.67	0.08	[...]	92.74
4	1.69	0.02	43.23	0.03	[...]	55.04
5	1.05	0.00	26.95	0.15	[...]	71.83
6	0.52	0.07	5.13	0.32	[...]	93.97
7	0.30	0.02	0.18	0.00	[...]	99.50

CPU



```
$ time ./encrypt.sh [...]  
[...]  
Encryption took 452.47 s
```

```
real      7m32.730s  
user      0m16.597s  
sys       7m17.231s
```

```
$ mpstat -P ALL 60
```

CPU	%usr	%nice	%sys	%iowait	[...]	%idle
all	0.68	0.02	12.41	0.15	[...]	86.73
0	0.32	0.02	1.39	0.60	[...]	97.68
1	0.25	0.02	0.13	0.02	[...]	99.58
2	0.82	0.03	15.59	0.03	[...]	83.53
3	0.50	0.00	6.67	0.08	[...]	92.74
4	1.69	0.02	43.23	0.03	[...]	55.04
5	1.05	0.00	26.95	0.15	[...]	71.83
6	0.52	0.07	5.13	0.32	[...]	93.97
7	0.30	0.02	0.18	0.00	[...]	99.50

CPU



```
$ time ./encrypt.sh [...]  
[...]  
Encryption took 452.47 s
```

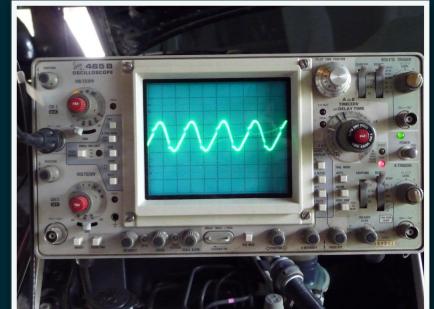
real	7m32.730s
user	0m16.597s
sys	7m17.231s

```
$ mpstat -P ALL 60  
CPU %usr %nice %sys %iowait [...] %idle  
all 0.68 0.02 12.41 0.15 [...] 86.73  
0 0.32 0.02 1.39 0.60 [...] 97.68  
1 0.25 0.02 0.13 0.02 [...] 99.58
```

2 * I/O: kein Problem
3 * Single Threaded (12.5% von 8 Cores = 1 Core)
4 * 96% Kernel Time?

```
6 0.52 0.07 5.13 0.32 [...] 93.9 /  
7 0.30 0.02 0.18 0.00 [...] 99.50
```

sys == Syscall



Messen z.B. mit systemtap *

```
#!/usr/bin/env stap
#
# This script lists the top 20 systemcalls in total
# start with stap -c command / stap -x pid

global syscalls_count

probe syscall.* {
    if (target() == pid()) {
        syscalls_count[name] <<< 1
    }
}

function print_systop () {
    printf ("%25s %10s\n", "SYSCALL", "COUNT")
    foreach (syscall in syscalls_count- limit 20) {
        printf("%25s %10d\n", syscall, @count(syscalls_count[syscall]))
    }
}

probe end {
    print_systop ()
}
```

* <https://sourceware.org/systemtap/>

Syscalls?



~125,000 write calls/s

bei einem Durchsatz von 141 kB/s

```
$ iostat -x -p dm-3 60
```

[...]	r/s	w/s	rkB/s	wkB/s	[...]	%util
[...]	0.17	6.50	0.67	141.07	[...]	0.47

Syscalls?



~125,000 write calls/s

bei einem Durchsatz von 141 kB/s

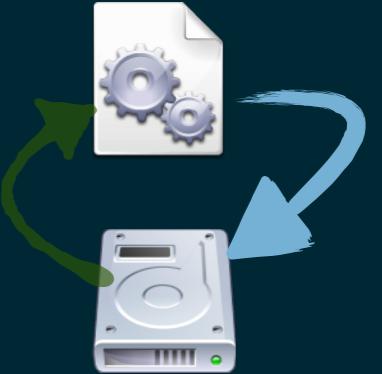
```
$ iostat -x -p dm-3 60
```

[...]	r/s	w/s	rkB/s	wkB/s	[...]	%util
[...]	0.17	6.50	0.67	141.07	[...]	0.47

- * I/O: kein Problem
- * 96% Kernel Time?
 - * 125k system calls für 141 kB!



write calls?



```
final OutputStream outputStream = new FileOutputStream(destFile);  
pgp.encryptAndSign(new FileInputStream(sourceFile), outputStream);
```



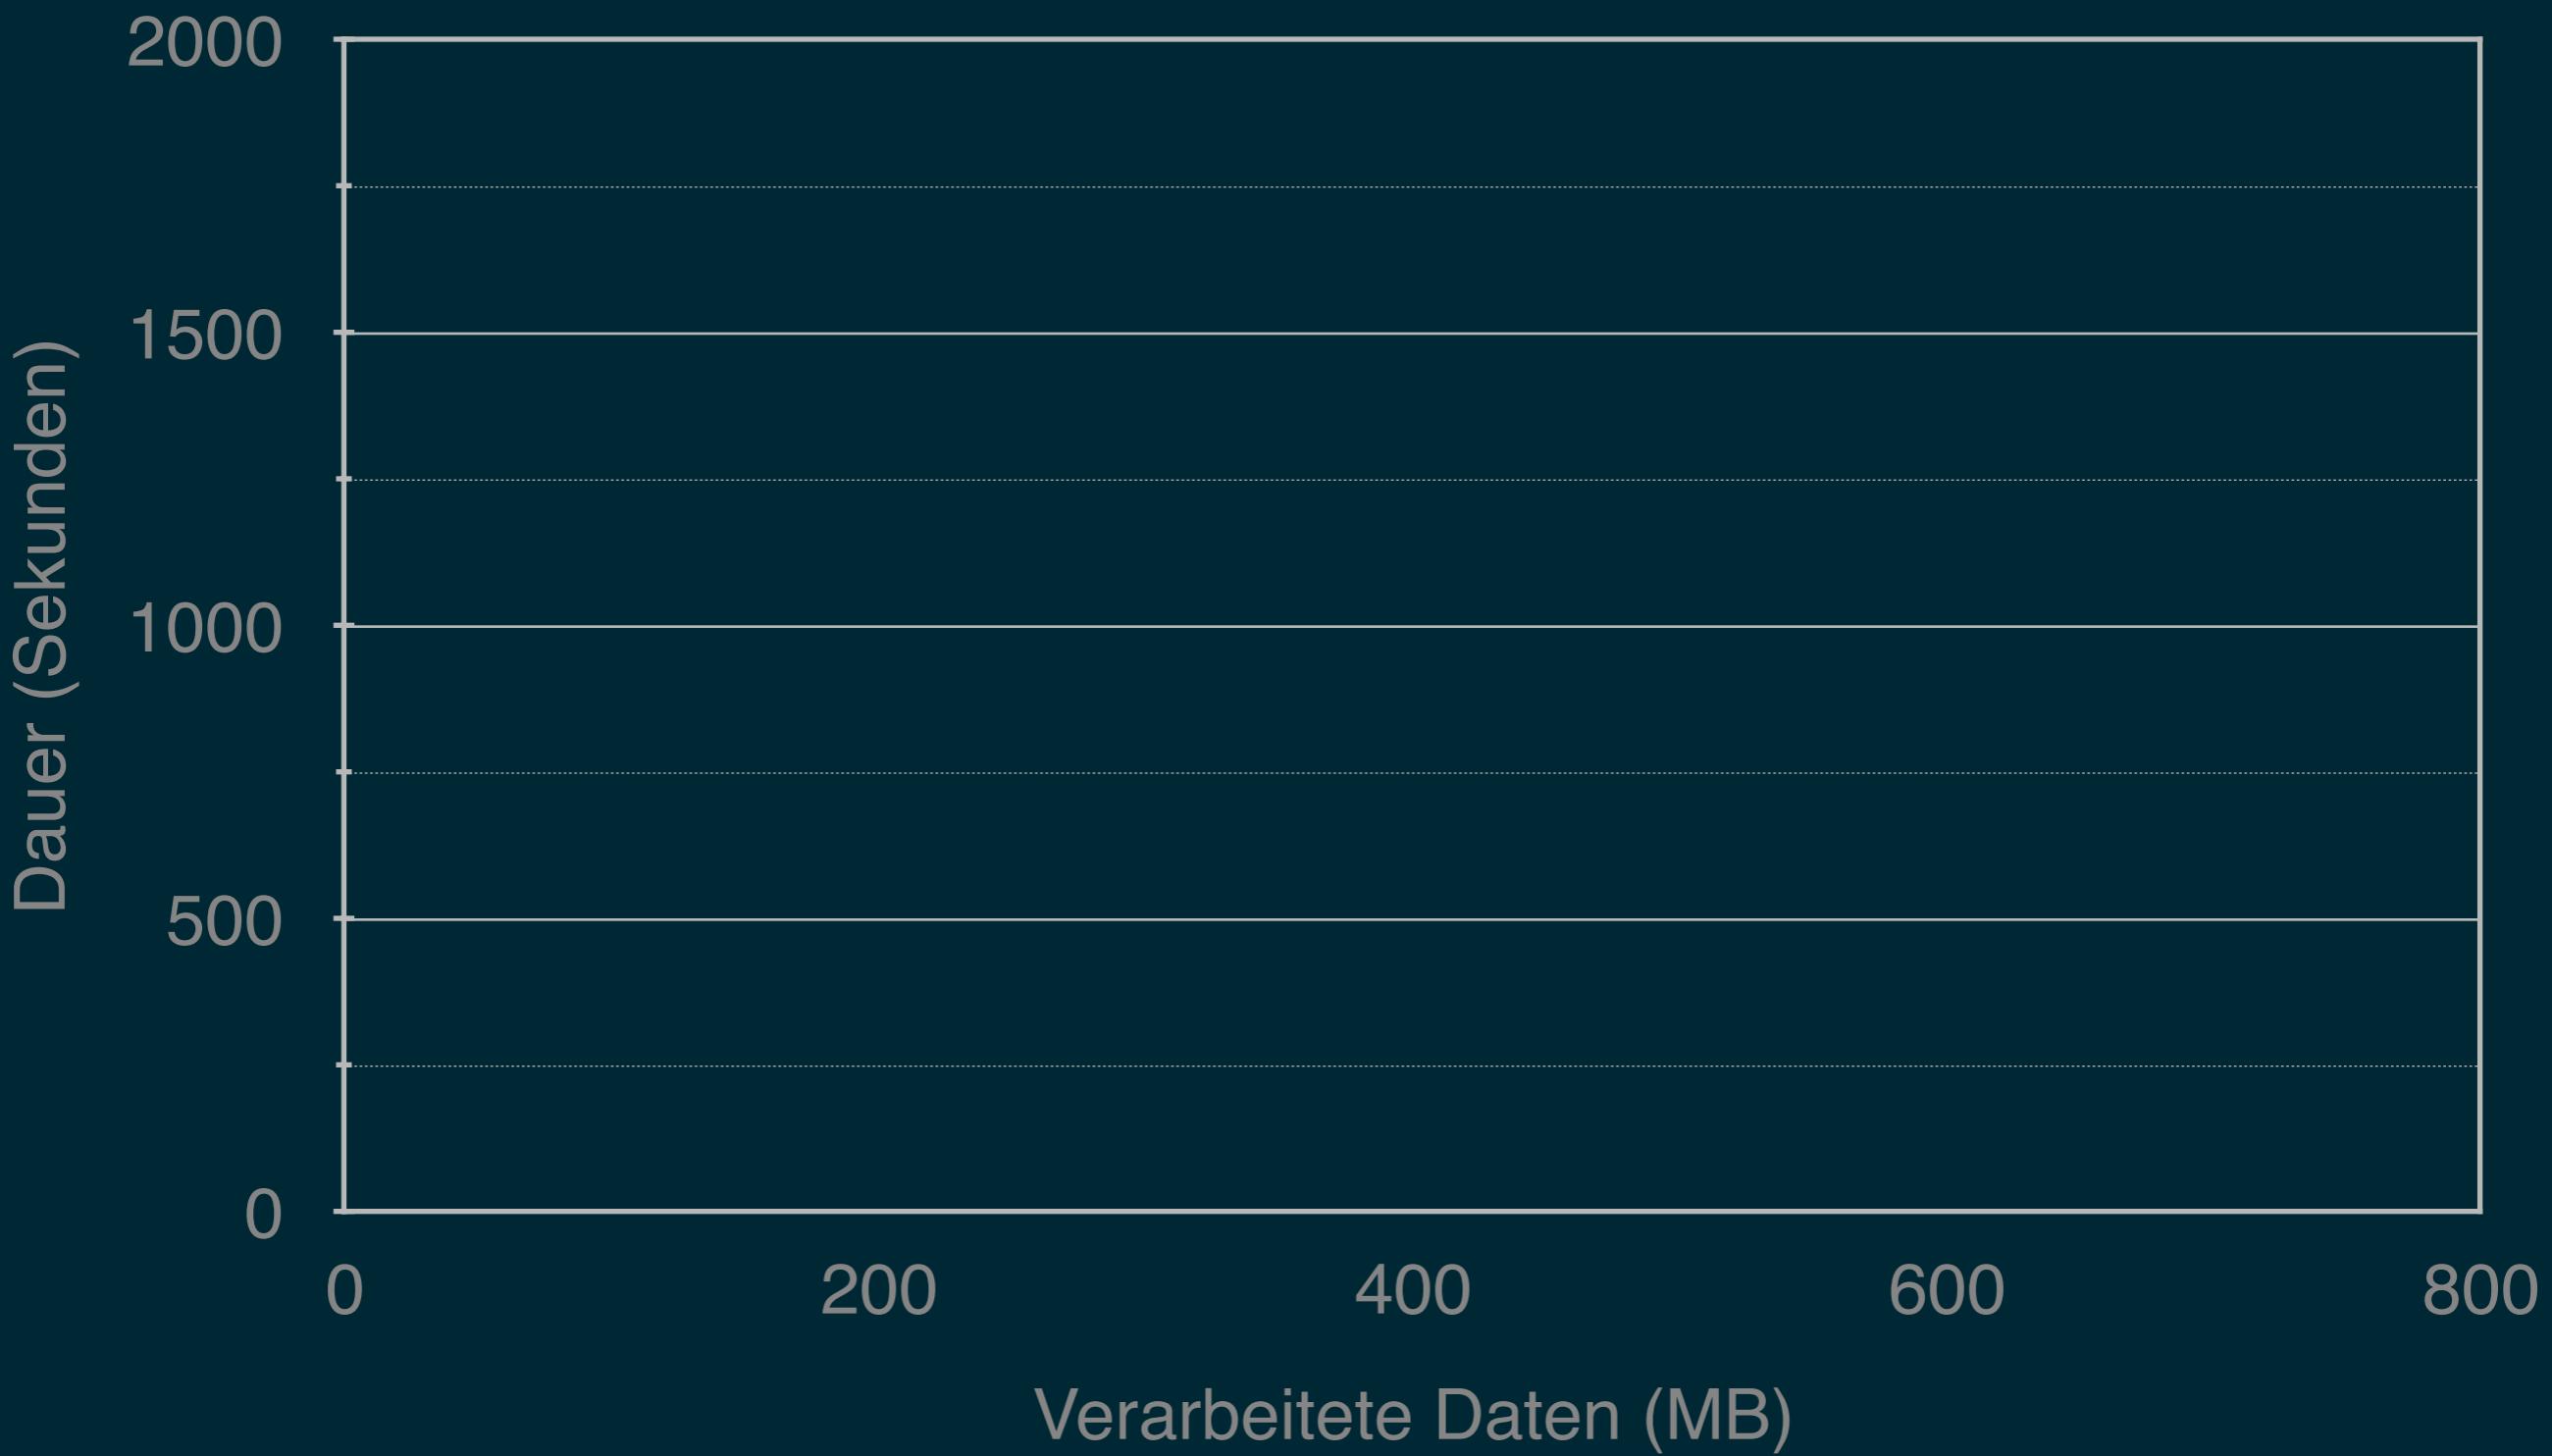
```
final OutputStream outputStream = new FileOutputStream(destFile);  
final OutputStream bufferedOutputStream = new BufferedOutputStream(outputStream, 8192);  
  
pgp.encryptAndSign(new FileInputStream(sourceFile), bufferedOutputStream);
```

Lastverhalten

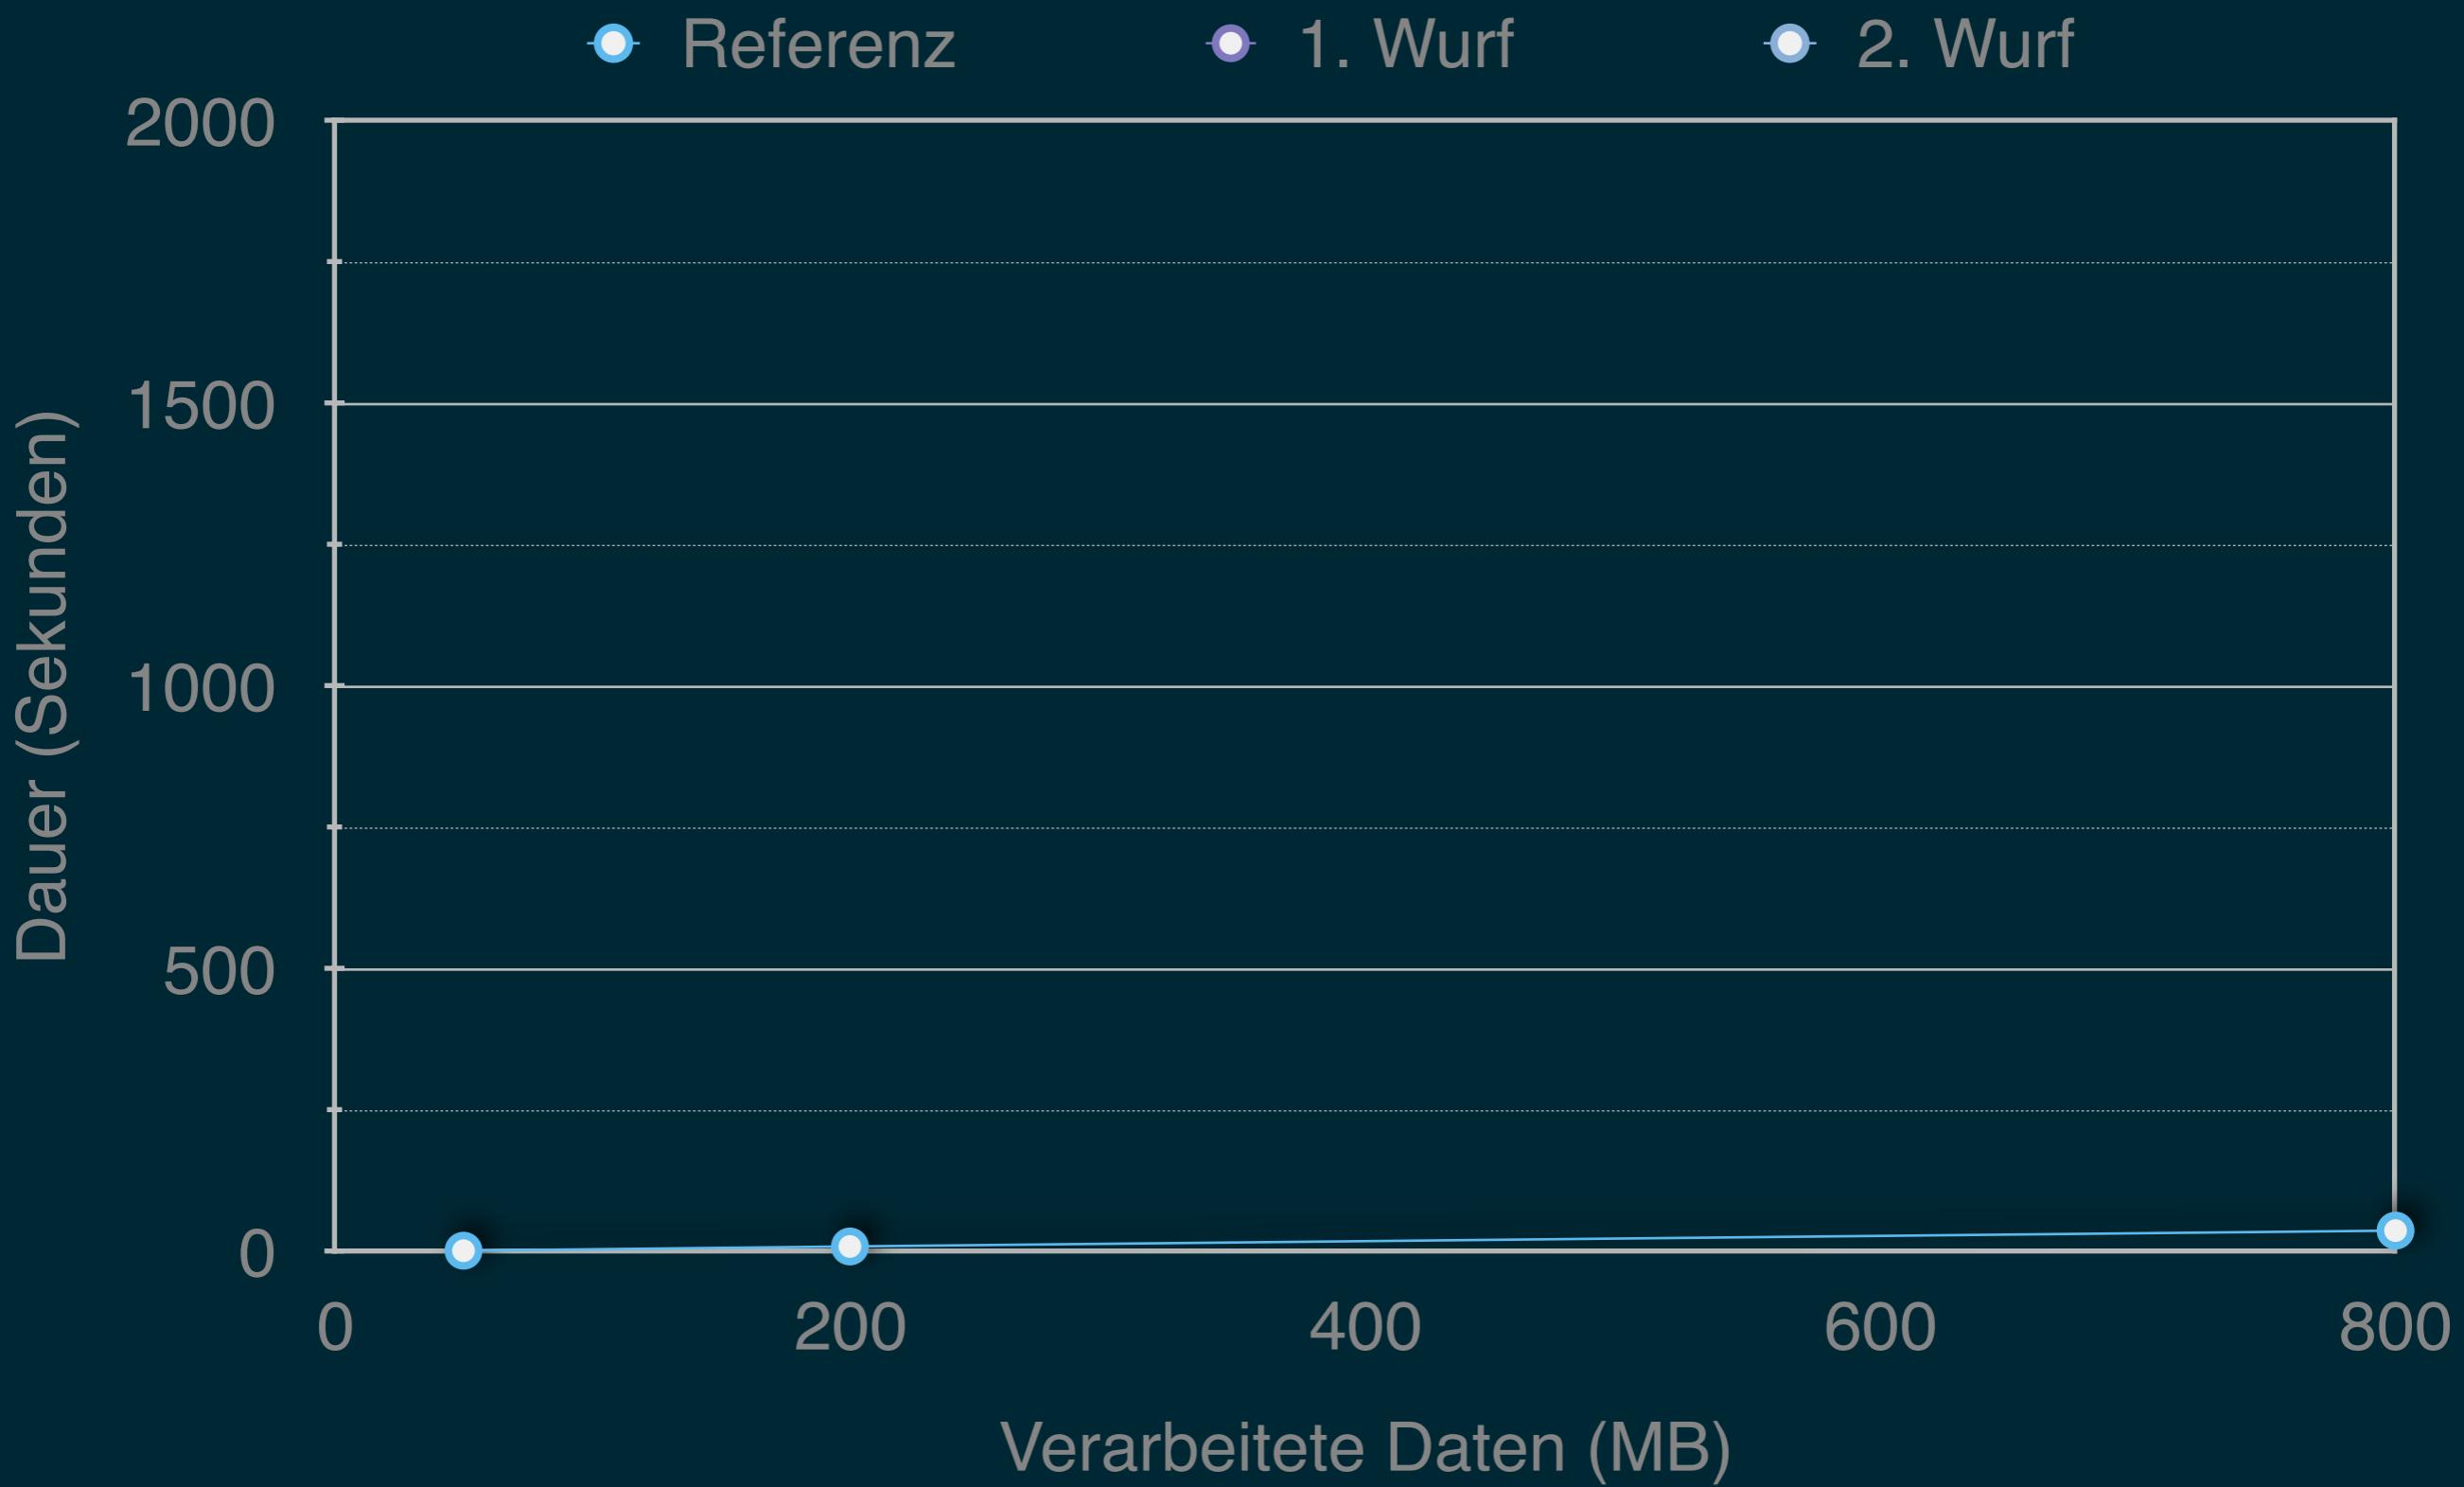
Wiederholung

Lastverhalten

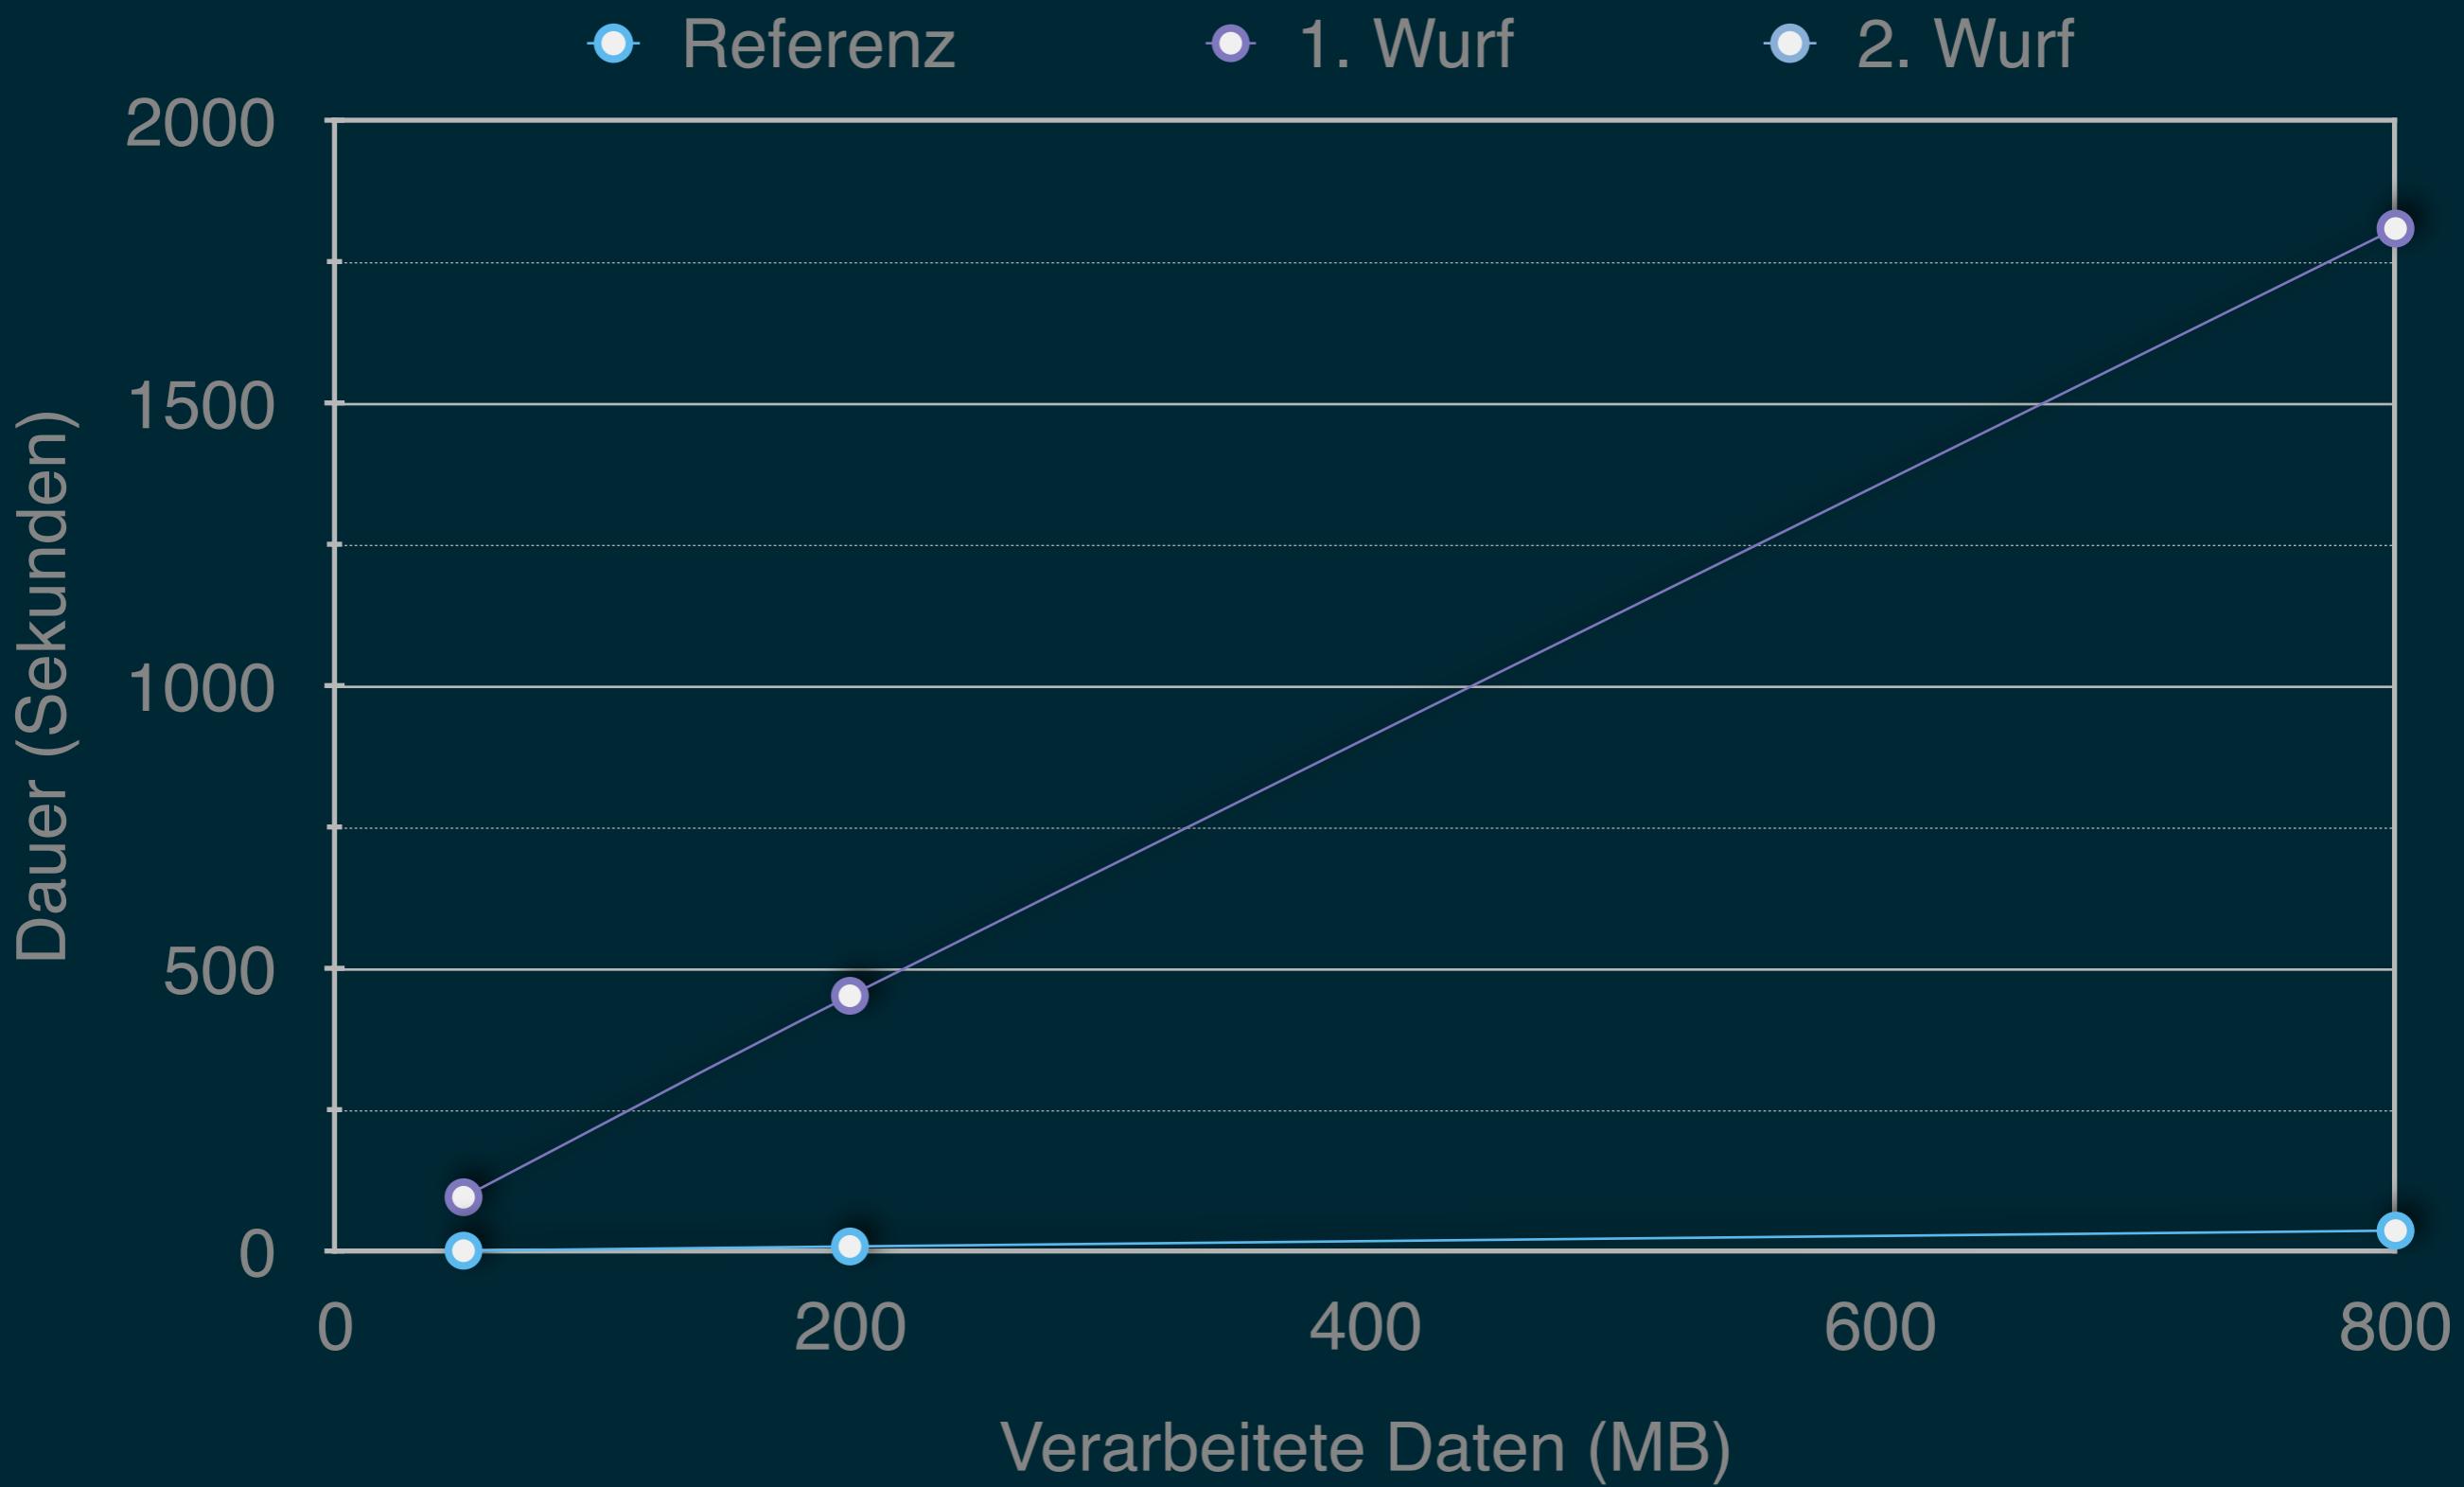
● Referenz ● 1. Wurf ● 2. Wurf



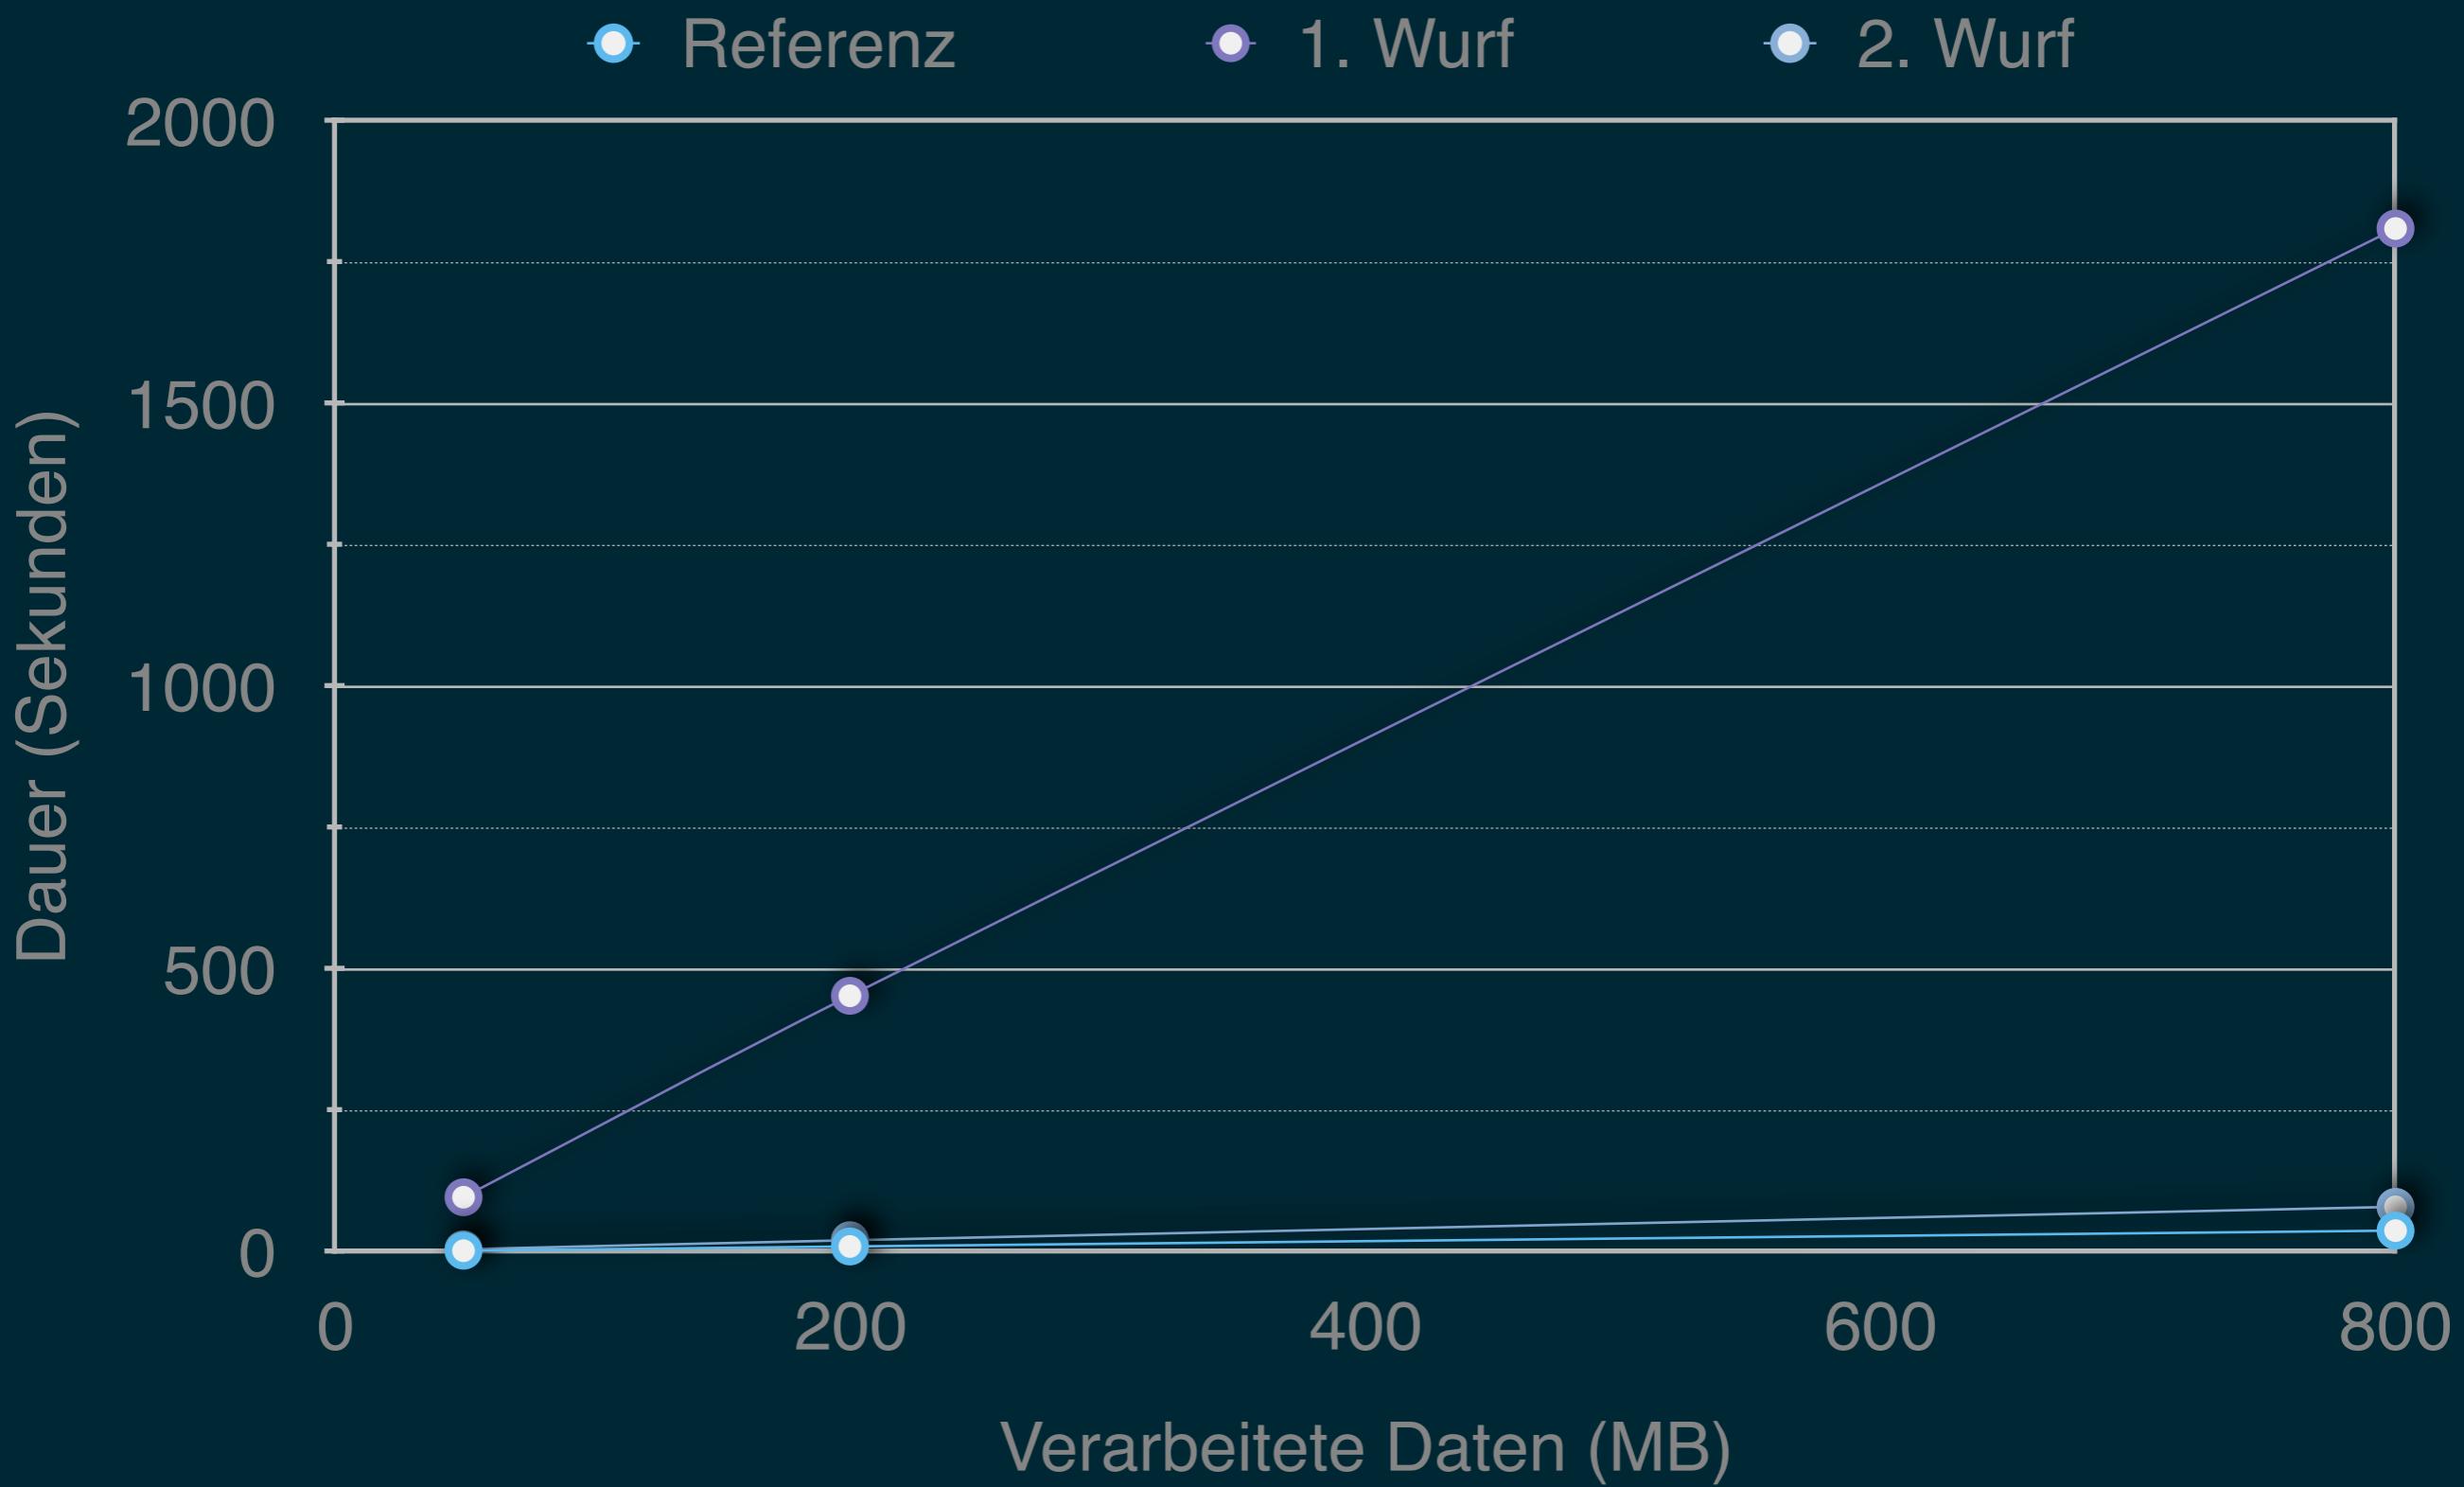
Lastverhalten



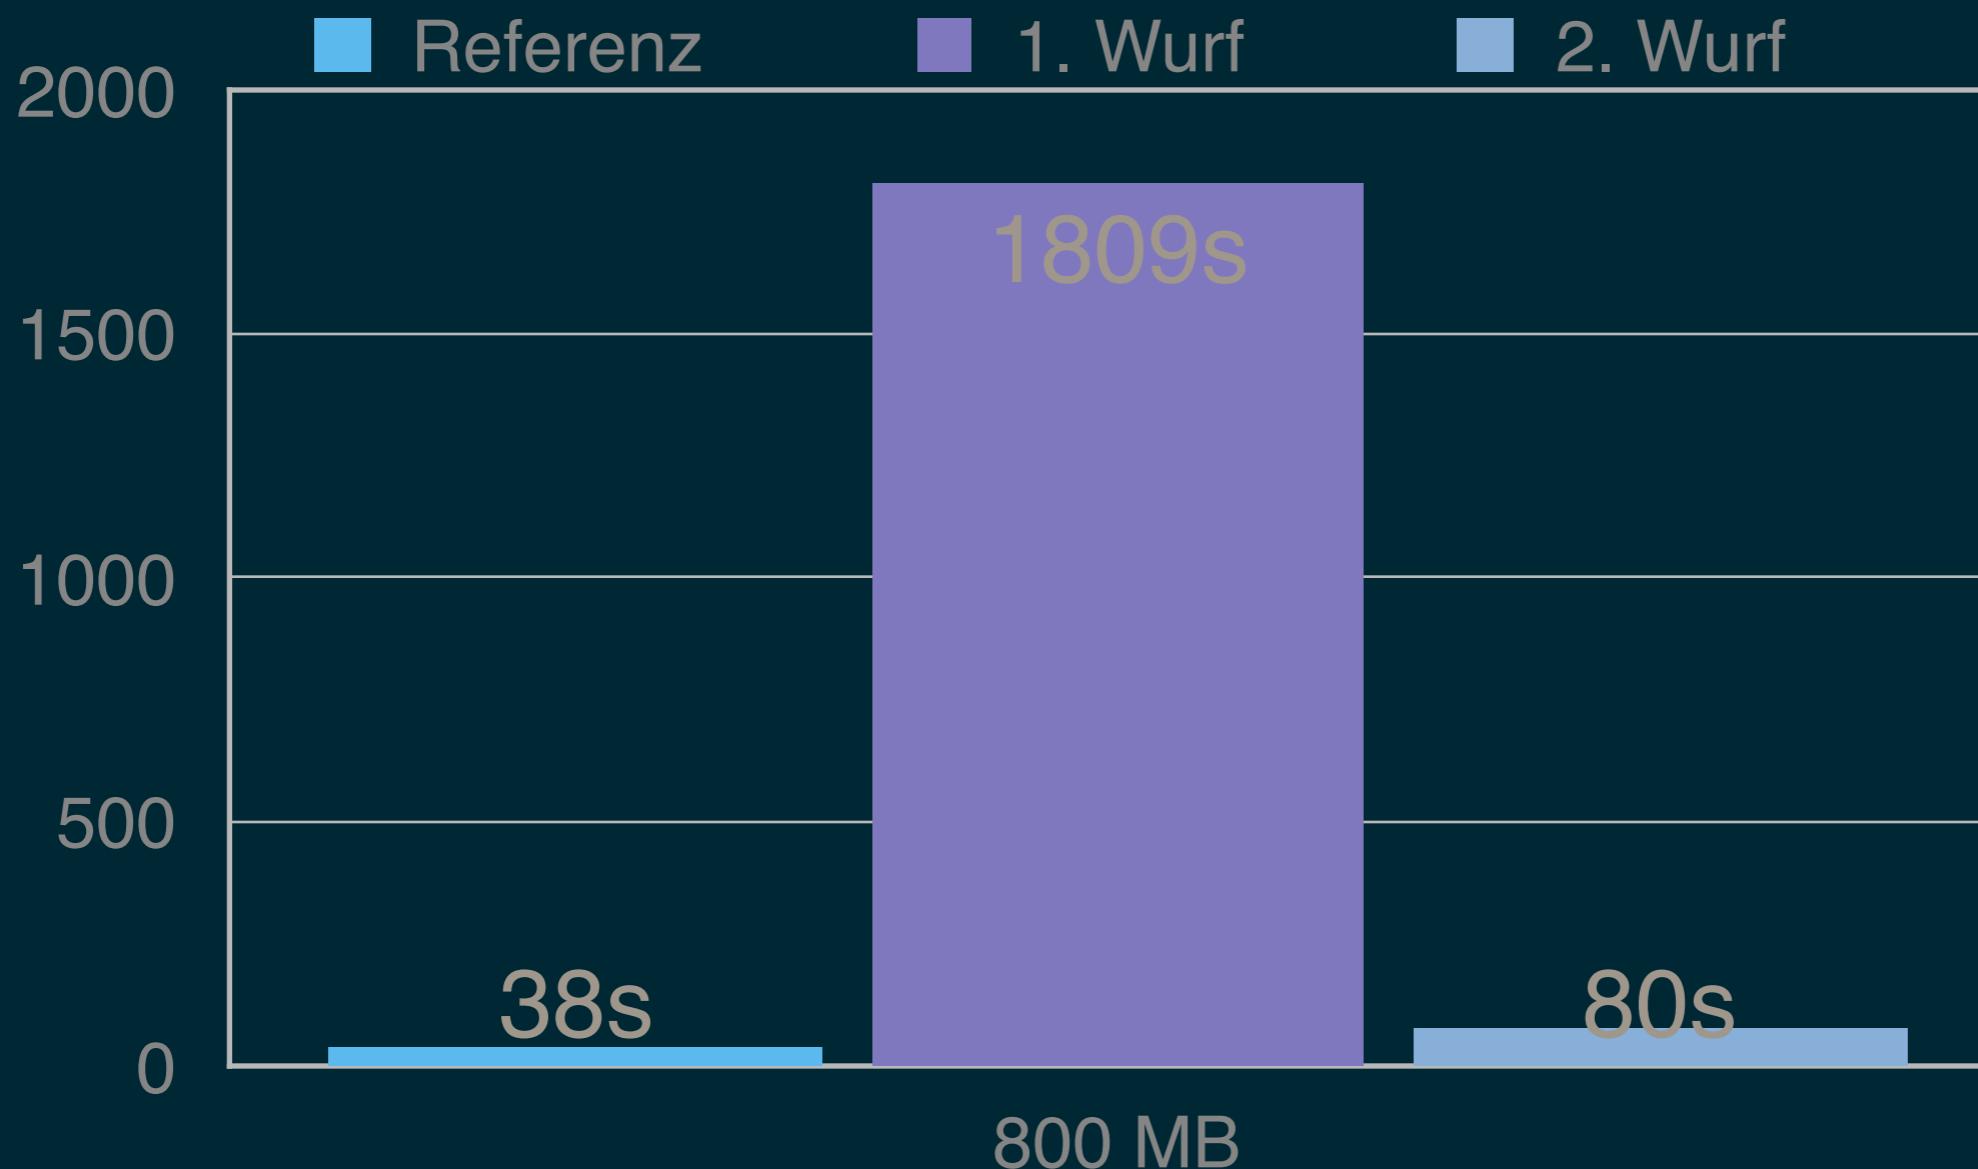
Lastverhalten



Lastverhalten



Lastverhalten



~ Faktor 47 zwischen Referenz und 1. Wurf
~ Faktor 2 zwischen Referenz und 2. Wurf





https://images.unsplash.com/40/x6YzbWWRq2sRhAacMjnl_Bangkok%20Indra%20market.jpg?ixlib=rb-0.3.5&q=80&fm=jpg&s=a3d1f05b8862f07db7a569c1bf1ecbc1 (Creative Commons Zero)



<https://images.unsplash.com/photo-1433616174899-f847df236857?ixlib=rb-0.3.5&q=80&fm=jpg&s=c1e8fe89e406303598e6d449f9b685eb> (Creative Commons Zero)



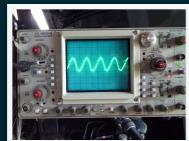
https://en.wikipedia.org/wiki/Hard_disk_drive#/media/File:Hdd_icon.svg (LGPL)



<https://www.dropbox.com/s/dyf59rjfhp09jaw/78H.jpg?dl=1> (Creative Commons Zero)



https://commons.wikimedia.org/wiki/File:Singapore_Road_Signs_-_Warning_Sign_-_Road_Narrows_On_Right.png (CC-BY 3)



https://upload.wikimedia.org/wikipedia/commons/3/39/RAV4_Tach_750rpm.JPG (public domain)



http://www.gratisography.com/pictures/77_1.jpg (Creative Commons Zero)